



Software Exploitation Basics

Mike Ham

You're here to learn software exploitation. What do YOU think it is, and what can successful exploitation result in?

Something More Relatable

- Halo 2 was a game that I played probably too much of
- There were many flaws in the software that allowed players to do some cool stuff, useful or not
- One particular flaw was called the “Super Bounce”
 - <https://www.youtube.com/watch?v=6gyZ9Hfd5IE>
- The flaw in the software was taken advantage of for some gain...exploits

Software Exploitation

- It takes on many forms and meanings so there isn't a *great* definition
- Taking advantage of a vulnerability in some software to get the system to execute code you provide
- Vulnerabilities are flaws or weaknesses in a system that can be *exploited*
 - Bad programmer, functions in software were designed poorly, etc.
- We exploit software for many purposes, it's up to you what the outcome should be!

Precursor

- If you're going to exploit something, you REALLY need to understand how it works or is designed to work in the first place
- Computers are super complex, different languages (C, Java, etc.) respond and operate differently, processors interpret commands differently, etc.
- Windows and other OS provide some security mechanisms to protect vulnerable programs from being exploited (DEP, ASLR, EMET)
- High value software is pretty big, may have millions of lines of code, well protected
- The list goes on....

Don't panic!

Why write exploits when I can use tool X or Y?

Exploitation Tools/Frameworks

- There are tons of these out there, Metasploit is a really common one
 - Armitage/Cobalt Strike, CORE Impact, Nmap
 - There absolutely is a time/place for using them, I'd actually encourage it
 - Someone had to write the exploits in those tools
- Anyone can fire up a tool and click “GO” - skiddies
 - If you want to make \$\$\$ and be 1337, learn to write some on your own or at least understand what you’re running

Machine & Compiler Setup

- If you are running a x64 Kali install, it will be easier to look at 32-bit programs
- Install the following software:
 - `apt-get install libc6-dev-i386 -y`
 - `apt-get install hexedit -y`
- Address Space Layout Randomization (ASLR), disable in terminal temporarily:
 - `setarch `uname -m` -R /bin/bash`

Software, Processors, RAM, and HDD

Software

- We write software in a high level language like C, C++, C#, Java, etc.
 - It's (pretty) easy for us to understand and read, but yet still efficient
 - Called **source code**
- Computers operate in **binary**, 1's and 0's to represent ALL information
 - They can't simply run a C program as is without compiling it
- Compiling takes the language we write in and translates it to a set of instructions that the processor can understand

C → Binary

- When you compile a C program, it gets turned into processor instructions
- Processors read the instructions in binary
- Obviously, binary is really tough to read
 - Use `xxd -b <file>` to read in binary

```
hello_world.c *
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 int main(void)
5 {
6     printf("Hello, World!");
7     exit(0);
8 }
```

```
0000f4e: 00000000 00000000 01010101 01001000 10001001 11100101 ..UH..
0000f54: 01001000 10000011 11101100 00010000 01001000 10001101 H...H.
0000f5a: 00111101 01001001 00000000 00000000 00000000 11000111 =I....
0000f60: 01000101 11111100 00000000 00000000 00000000 00000000 E.....
0000f66: 10110000 00000000 11101000 00010001 00000000 00000000 .....
0000f6c: 00000000 00110001 11111111 10001001 01000101 11111000 .1..E.
0000f72: 11101000 00000001 00000000 00000000 00000000 10010000 .....
0000f78: 11111111 00100101 10010010 00000000 00000000 00000000 .%...
0000f7e: 11111111 00100101 10010100 00000000 00000000 00000000 .%...
0000f84: 01001100 10001101 00011101 01111101 00000000 00000000 L..}..
0000f8a: 00000000 01000001 01010011 11111111 00100101 01101101 .AS.%m
0000f90: 00000000 00000000 00000000 10010000 01101000 00000000 ....h.
0000f96: 00000000 00000000 00000000 11101001 11100110 11111111 .....
0000f9c: 11111111 11111111 01101000 00001100 00000000 00000000 ..h...
0000fa2: 00000000 11101001 11011100 11111111 11111111 11111111 .....
0000fa8: 01001000 01100101 01101100 01101100 01101111 00101100 Hello,
0000fae: 00100000 01010111 01101111 01110010 01101100 01100100 World
0000fb4: 00100001 00000000 00000000 00000000 00000001 00000000 !.....
0000fba: 00000000 00000000 00011100 00000000 00000000 00000000 .....
0000fc0: 00000000 00000000 00000000 00000000 00011100 00000000 .....
0000fc6: 00000000 00000000 00000000 00000000 00000000 00000000 .....
0000fcc: 00011100 00000000 00000000 00000000 00000010 00000000 .....
0000fd2: 00000000 00000000 01010000 00001111 00000000 00000000 ..P...
0000fd8: 00110100 00000000 00000000 00000000 00110100 00000000 4...4.
0000fde: 00000000 00000000 01110000 00001111 00000000 00000000 ..x...
0000fe4: 00000000 00000000 00000000 00000000 00110100 00000000 ....4.
0000fea: 00000000 00000000 00000011 00000000 00000000 00000000 .....
```

Assembly

- We don't always have source code available, and we can't read binary
- There is a language in-between called **assembly**
 - It's not super easy to read, but we can see the processor instructions and work from those - takes practice!
- There is a tool called **objdump** to view assembly in Kali

```
0804842b <main>:  
 804842b: 8d 4c 24 04      lea    0x4(%esp),%ecx  
 804842f: 83 e4 f0      and    $0xffffffff,%esp  
 8048432: ff 71 fc      pushl  -0x4(%ecx)  
 8048435: 55              push   %ebp  
 8048436: 89 e5          mov    %esp,%ebp  
 8048438: 51              push   %ecx  
 8048439: 83 ec 04      sub    $0x4,%esp  
 804843c: 83 ec 0c      sub    $0xc,%esp  
 804843f: 68 f0 84 04 08  push   $0x80484f0  
 8048444: e8 a7 fe ff ff  call   80482f0 <printf@plt>  
 8048449: 83 c4 10      add    $0x10,%esp  
 804844c: 83 ec 0c      sub    $0xc,%esp  
 804844f: 6a 00          push   $0x0  
 8048451: e8 ba fe ff ff  call   8048310 <exit@plt>  
 8048456: 66 90          xchg   %ax,%ax  
 8048458: 66 90          xchg   %ax,%ax  
 804845a: 66 90          xchg   %ax,%ax  
 804845c: 66 90          xchg   %ax,%ax  
 804845e: 66 90          xchg   %ax,%ax
```

How Software is Loaded

- Software is stored on your HDD in binary and exists permanently (Word, Steam)
- When you want to run software, a copy of it gets moved from the HDD over into your computers RAM
 - Why?
- The processor then can read the instructions within the software and run the program
- When you close the software, RAM is cleared out and freed for something else

RAM/Memory

- Memory doesn't just store programs
 - It also is capable of storing data (numbers, words, etc.) - variables
- CPU can jump to any line or portion of RAM in its operations
- CPU can access one row of RAM just as easily as any other
 - (Random Access Memory)

Ram is Like a Spreadsheet

Address

Memory - C:\Users\malware\Desktop\calc.exe - WinDbg:10.0.10586.567 X86											
Virtual: 0x`76f069ef				Previous		Display format:		Bit		Next	
76f068f5	10010100	11111001	01110110	00110011	11000000	01011110	01000000	01011011			
76f068fd	10001011	11100101	01011101	11000010	00001100	00000000	01010000	11101000			
76f06905	00110111	01001111	00000001	00000000	11001100	11001100	11001100	11001100			
76f0690d	11001100	11001100	11001100	10001011	11111111	01010101	10001011	11101100			
76f06915	00001111	10110111	01000101	00001100	01100110	10000101	11000000	00001111			
76f0691d	10000100	11000010	00000000	00000000	00000000	10001011	01010101	00001000			
76f06925	11110110	11000010	00000001	00001111	10000101	10110110	00000000	00000000			
76f0692d	00000000	01100110	00111011	01000010	00110100	00001111	10000101	10101100			
76f06935	00000000	00000000	00000000	10000000	01111010	01001100	00000000	01010011			
76f0693d	10001010	01011101	00010000	01010110	10001011	01110101	00011000	01010111			
76f06945	10001011	01111101	00010100	01110100	00111001	10001010	01000010	01001101			
76f0694d	00111010	11011100	01110110	00000100	10000100	11000000	01110101	00101110			
76f06955	11110110	01000010	01001000	01000000	01110100	000000110	10001011	11000111			
76f0695d	00001011	11000110	01110100	01111000	10001011	01001010	01000000	10001011			
76f06965	01000010	01000100	00100011	11001111	00100011	110000110	00001011	11001000			
76f0696d	01110100	00010100	10001011	01000010	00111000	10001011	01001010	00111100			
76f06975	00100011	11000111	00100011	11001110	00111011	01000010	00111000	01110101			
76f0697d	00000101	00111011	01001010	00111100	01110100	01010110	10000000	10111010			
76f06985	11000100	00000000	00000000	00000000	00000000	01110100	01010001	10001010			
76f0698d	10000010	11000101	00000000	00000000	00000000	00111010	11011000	01110110			
76f06995	00000100	10000100	11000000	01110101	01000011	11110110	10000010	11000000			
76f0699d	00000000	00000000	00000000	01000000	01110100	000000110	10001011	11000111			
76f069a5	00001011	11000110	01110100	00110000	10001011	10001010	10111000	00000000			
76f069ad	00000000	00000000	10001011	10000010	10111100	00000000	00000000	00000000			
76f069b5	00100011	11001111	00100011	11000110	00001011	11001000	01110100	00100000			
76f069bd	10001011	10001010	10110000	00000000	00000000	00000000	10001011	11000001			
76f069c5	10001011	10011010	10110100	00000000	00000000	00000000	00100011	11000111			
76f069cd	10001011	11010011	00100011	11010110	00111011	11000001	01110101	00001000			
76f069d5	00111011	11010011	01110101	00000100	10110000	00000001	11101011	00000010			
76f069dd	00110010	11000000	01011111	01011110	01011011	11101011	00000010	00110010			
76f069e5	11000000	01011101	11000010	00010100	00000000	11001100	11001100	11001100			
76f069ed	11001100	11001100	11001100	10001011	11111111	01010101	10001011	11101100			

Very Quick Recap

C Source
Code

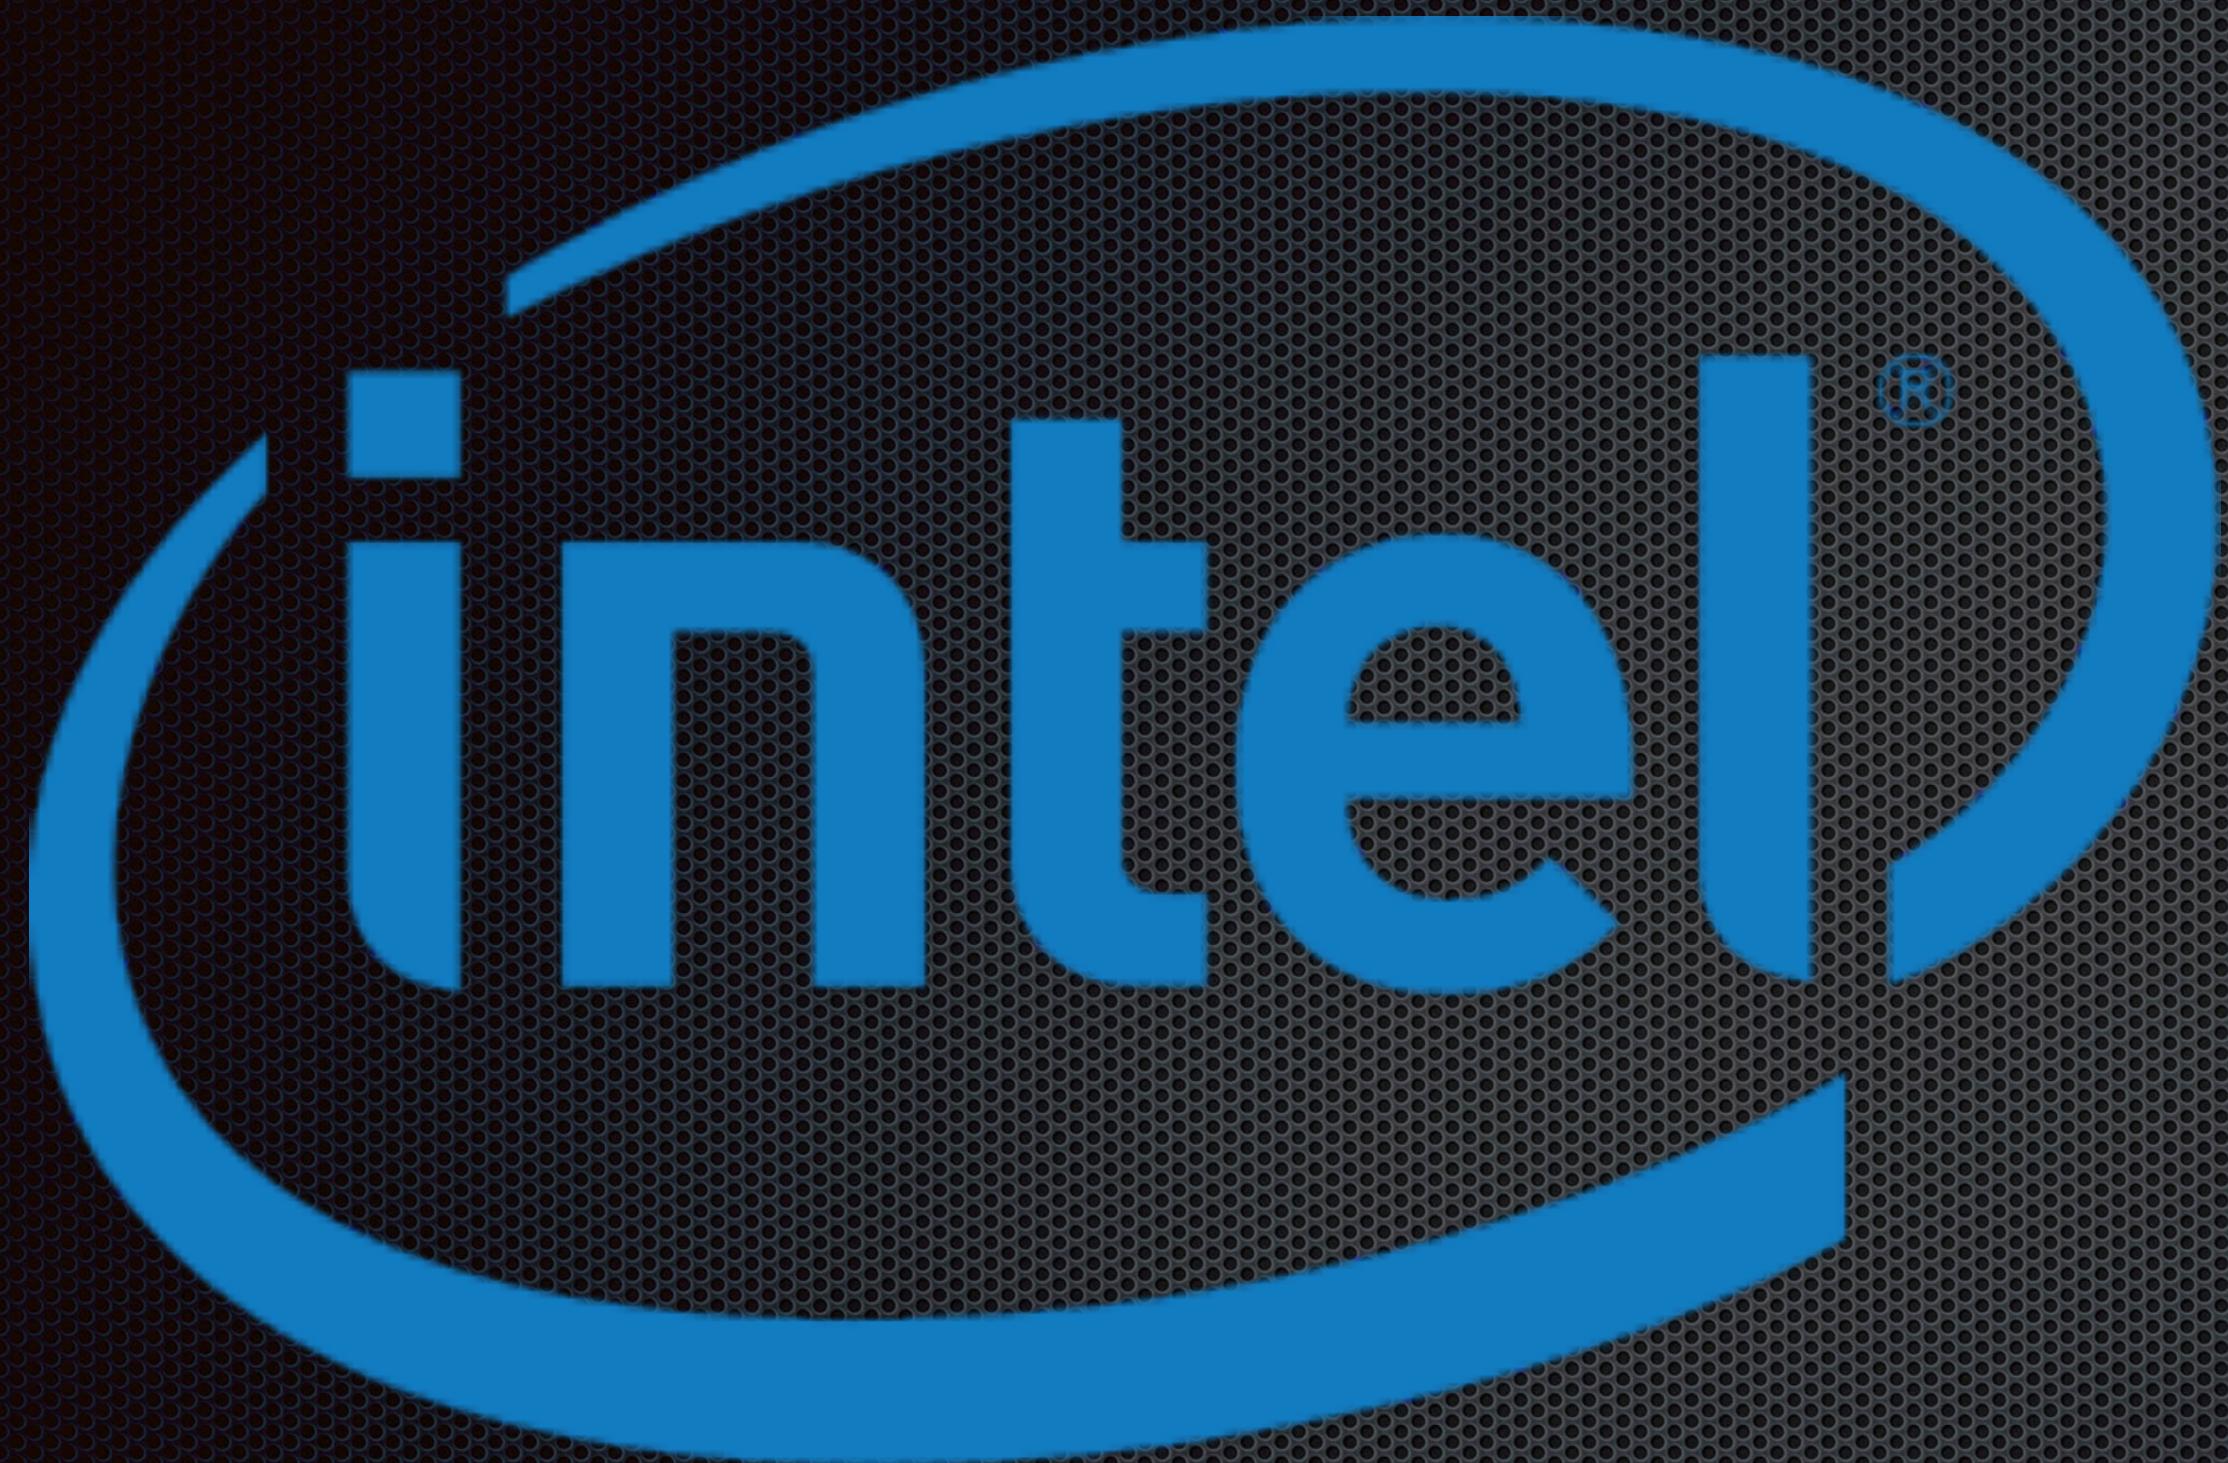
Compiler

Binary
(1's & 0's)

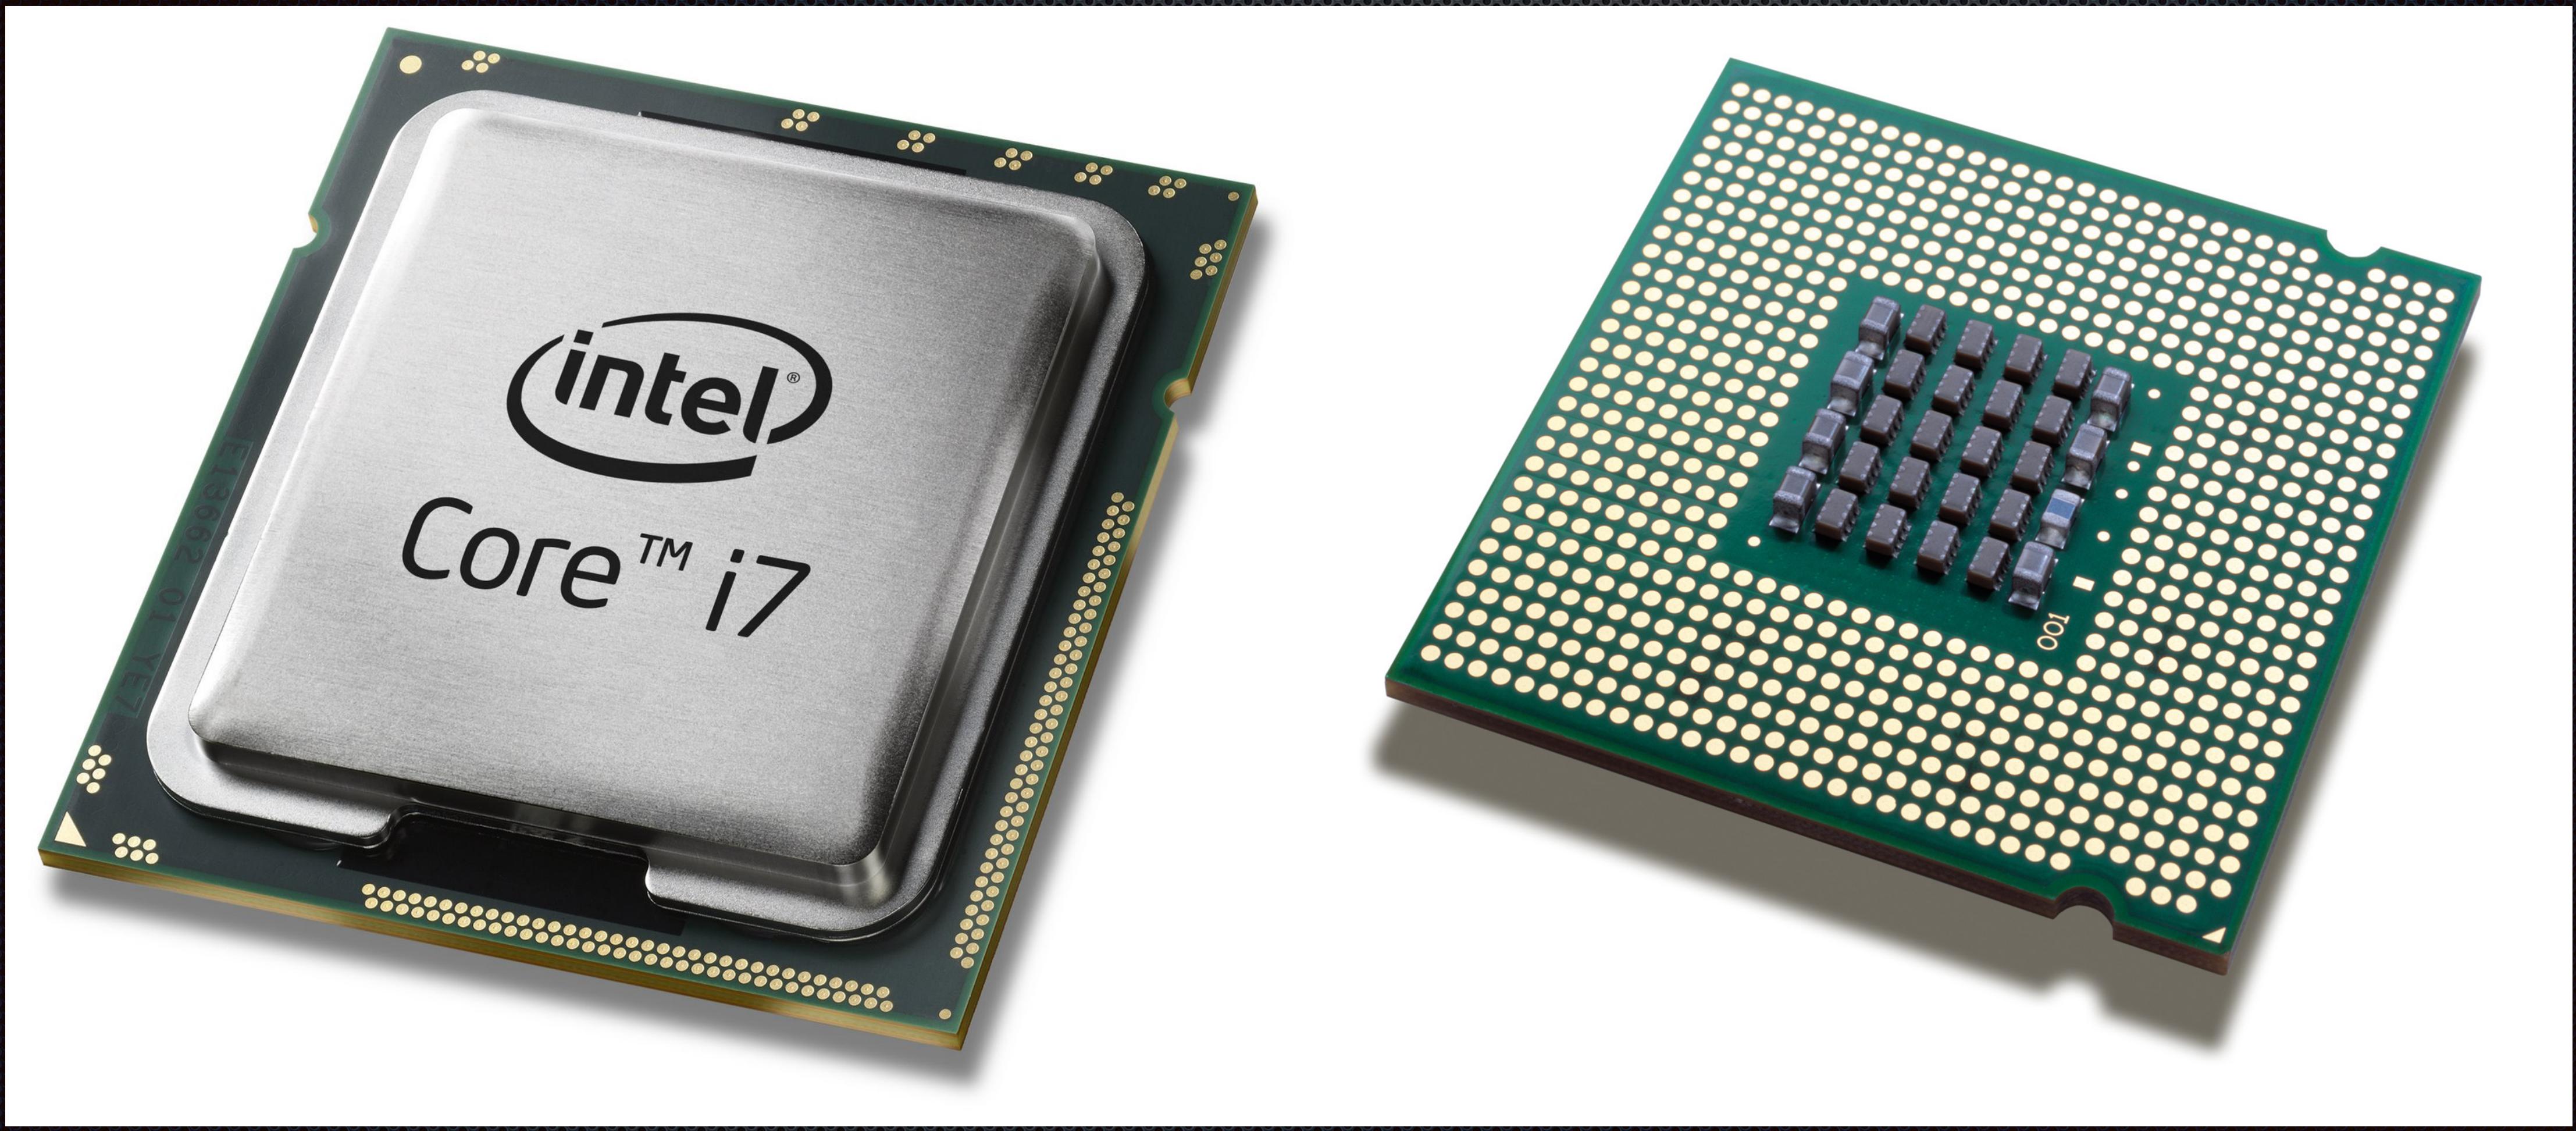
RAM

Ready to use!

The Processor



Physical Processor



More on Processors

- Processors come in different capacities and speeds
- Really common to find 32-bit and 64-bit in a computer
 - Effectively, how large of a number can the processor work with
- 32-bit binary number:
 - 11111111 11111111 11111111 11111111
- Kind of hard to read, isn't it?

Data Sizes

- Everything in a computer is really based off of binary
- 1 digit (1 or 0) is called a **bit**
- 8 bits put together is a **byte** (11010111)
 - Single letters (characters) or numbers are typically a byte in size
- 1024 bytes is a **kilobyte**
- 1024 kilobytes is a **megabyte**

32-Bit Processor

- They prefer to deal in 32-bit numbers, how many bytes are those numbers?
 - 4 bytes! ($32\text{bits}/8\text{bits} = 4 \text{ bytes}$)
- Often times, you will see your memory addresses in RAM represented by a 4 byte address
 - 0xaf12d34b <- random address I made up
 - It is represented in hexadecimal, just a little shorter/easier to read
- 2 hex digits (AF) = 1 byte

RAM - A Closer Look

The screenshot shows a memory dump from WinDbg for the process calc.exe at address 0x76f069ef. The display format is set to Bit. The memory dump consists of two columns of binary data. The first column is labeled 'Virtual' and the second is labeled 'Physical'. A blue box highlights a 4-byte block starting at address 76f069e5, which is labeled '4B'. Another blue box highlights a single byte at address 76f069e6, which is labeled '1B'.

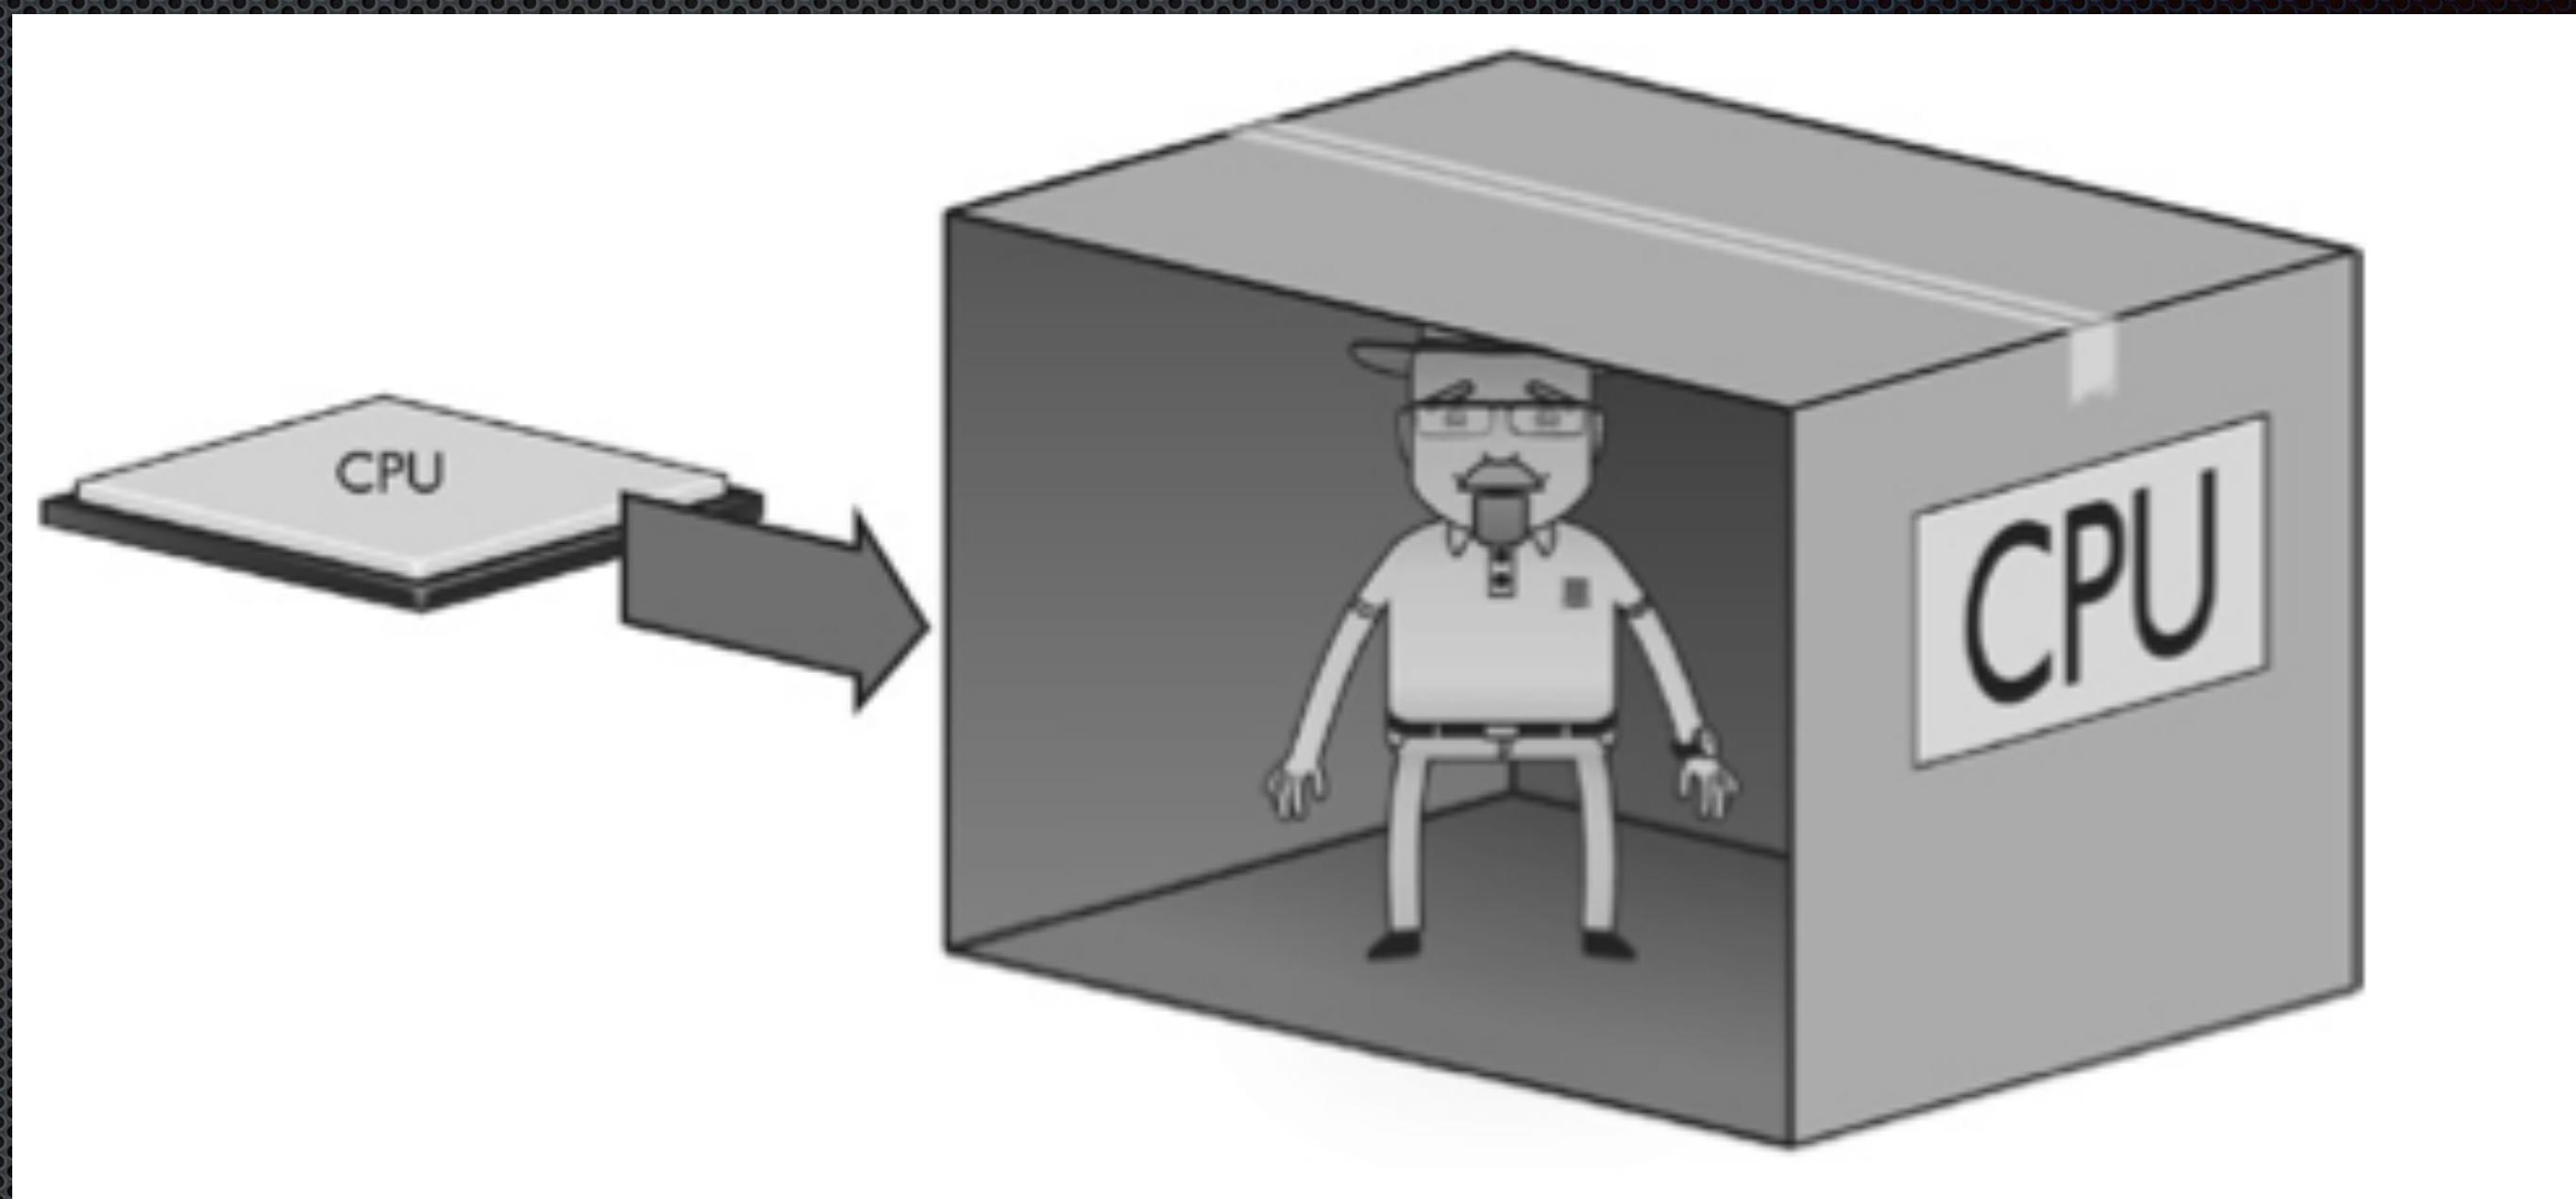
Virtual	Physical
76f068f5	10010100 11111000 01110110 00110011 11000000 01011110 01000000 01011011
76f068fd	10001011 11100100 01011101 11000010 00001100 00000000 01010000 11101000
76f06905	00110111 01001111 00000001 00000000 11001100 11001100 11001100 11001100
76f0690d	11001100 11001100 11001100 10001011 11111111 01010101 10001011 11101100
76f06915	00001111 10110111 01000101 00001100 01100110 10000101 11000000 00001111
76f0691d	10000100 11000010 00000000 00000000 10001011 01010101 00001000
76f06925	11101110 11000010 00000001 00001111 10000101 10110110 00000000 00000000
76f0692d	00000000 01100110 00111011 01000010 00110100 00001111 10000101 10101100
76f06935	00000000 00000000 00000000 10000000 01110100 01001100 00000000 01010011
76f0693d	10001010 01011101 00010000 01010110 10001011 01110101 00011000 01010111
76f06945	10001011 01111010 00010100 01110100 00111001 10001010 01000010 01001101
76f0694d	00111010 11011100 01110110 00000100 10000100 11000000 01110101 00101110
76f06955	11101110 01000010 01001000 01000000 01110100 00000110 10001011 11000111
76f0695d	00001011 11000110 01110100 01111000 10001011 01001010 01000000 10001011
76f06965	01000010 01000100 00100011 11001111 00100011 11000110 00001011 11001000
76f0696d	01110100 00010100 10001011 01000010 00111000 10001011 01001010 00111100
76f06975	00100011 11000111 00101111 11001110 00111011 01000010 00111000 01110101
76f0697d	00000101 00111011 01001010 00111100 01110100 01010110 10000000 10111010
76f06985	11000100 00000000 00000000 00000000 00000000 01110100 01010001 10001010
76f0698d	00000010 11000101 00000000 00000000 00000000 00111010 11011000 01110110
76f06995	00000100 10000100 11000000 01110101 01000011 11110110 10000010 11000000
76f0699d	00000000 00000000 00000000 01000000 01110100 00000110 10001011 11000111
76f069a5	00001011 11000111 01110100 00110000 10001011 10001010 10111000 00000000
76f069ad	00000000 00000000 10001011 10000010 10111100 00000000 00000000 00000000
76f069b5	00100011 11001111 00100011 11000110 00001011 11001000 01110100 00100000
76f069bd	10001011 10001011 10110000 00000000 00000000 10001011 11000001
76f069c5	10001011 10011010 10110100 00000000 00000000 00100011 11000111
76f069cd	10001011 11010011 00100011 11010110 00111011 11000001 01110101 00001000
76f069d5	00111011 11010011 01110101 00000100 10110000 00000001 11101011 00000010
76f069dd	00110010 11000000 01011111 01011110 01011011 11101011 00000010 00110010
76f069e5	11000000 01011101 11000010 00010100 00000000 11001100 11001100 11001100
76f069ed	11001100 11001100 11001100 10001011 11111111 01010101 10001011 11101100

From RAM to the Processor

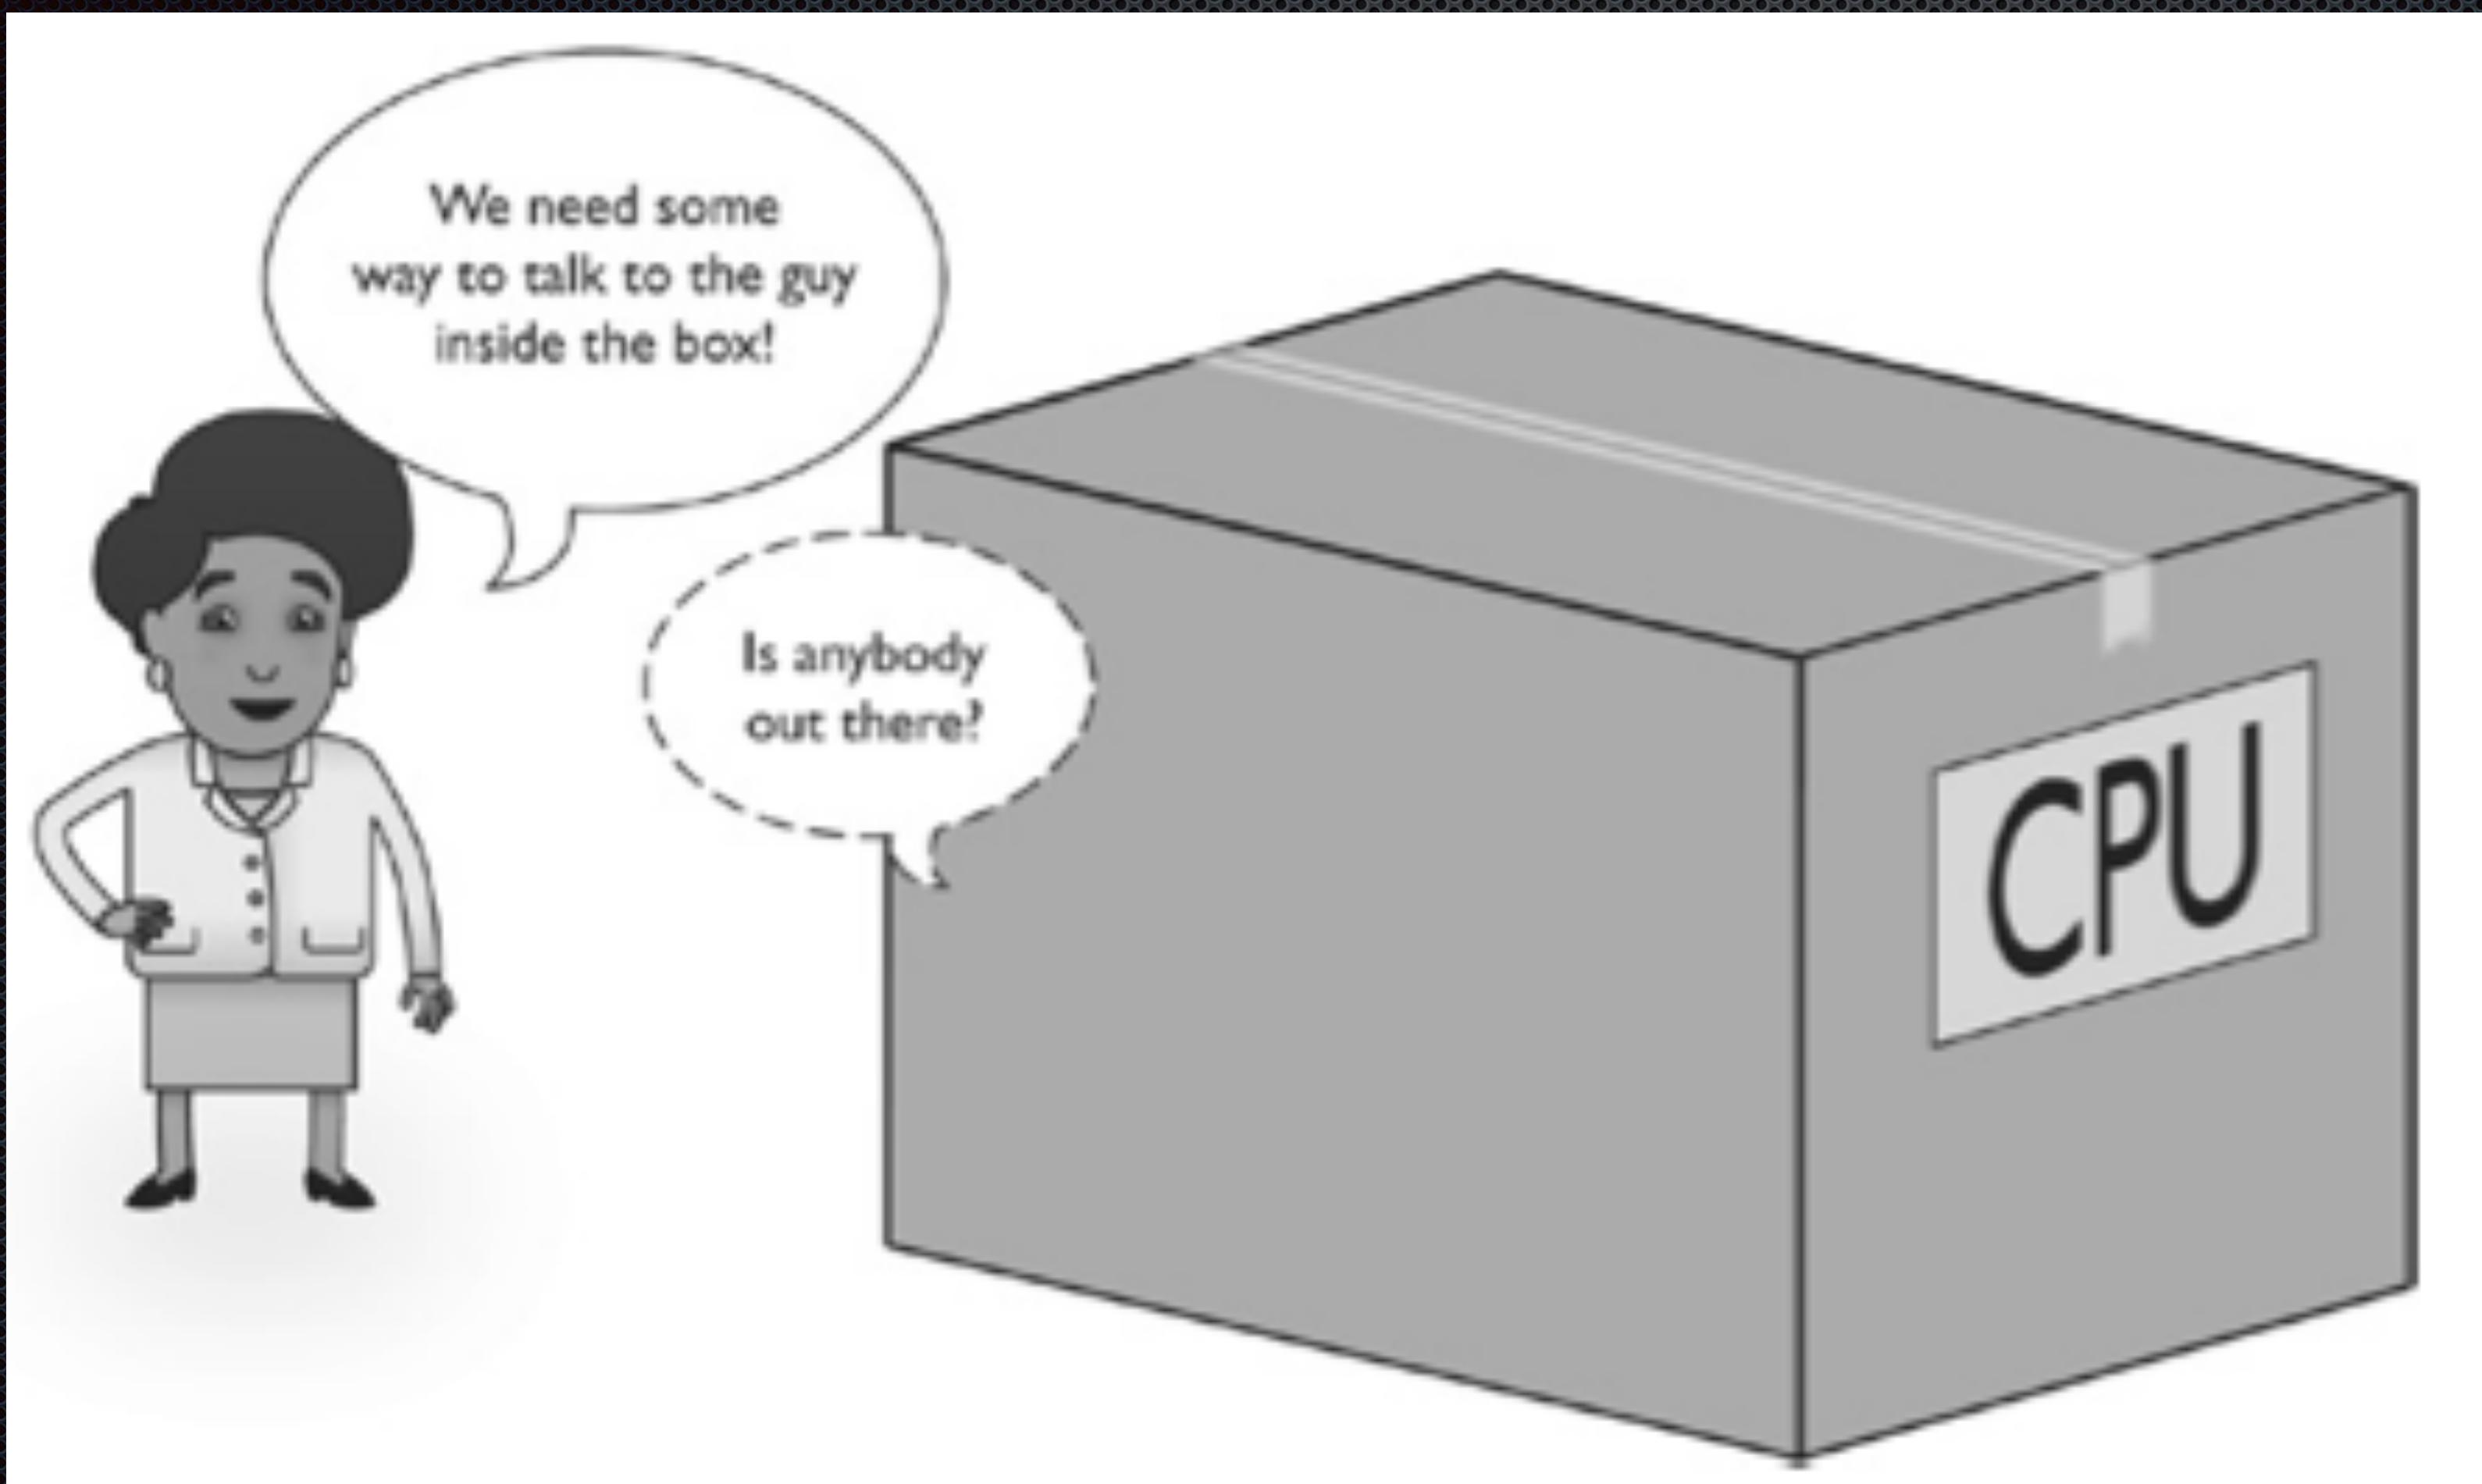
- Your program is currently sitting in RAM
- Processor needs to interpret the binary instructions (run the program)
- If we want to understand exploits, we need to understand a little more about the CPU/RAM and how they interact
- You can get really into the gritty details, but we'll keep it high level

A Man in a Box

- Think of him as a math genius
 - Can do nearly any equation
 - Provides answers really quick



Talking to the man in a box...



Let's Assume It's an 8-bit Processor

- Think of a set of 16 lightbulbs
 - 8 of them are inside the box
 - 8 of them are outside of the box
 - Each bulb has a switch that will turn on or off a pair

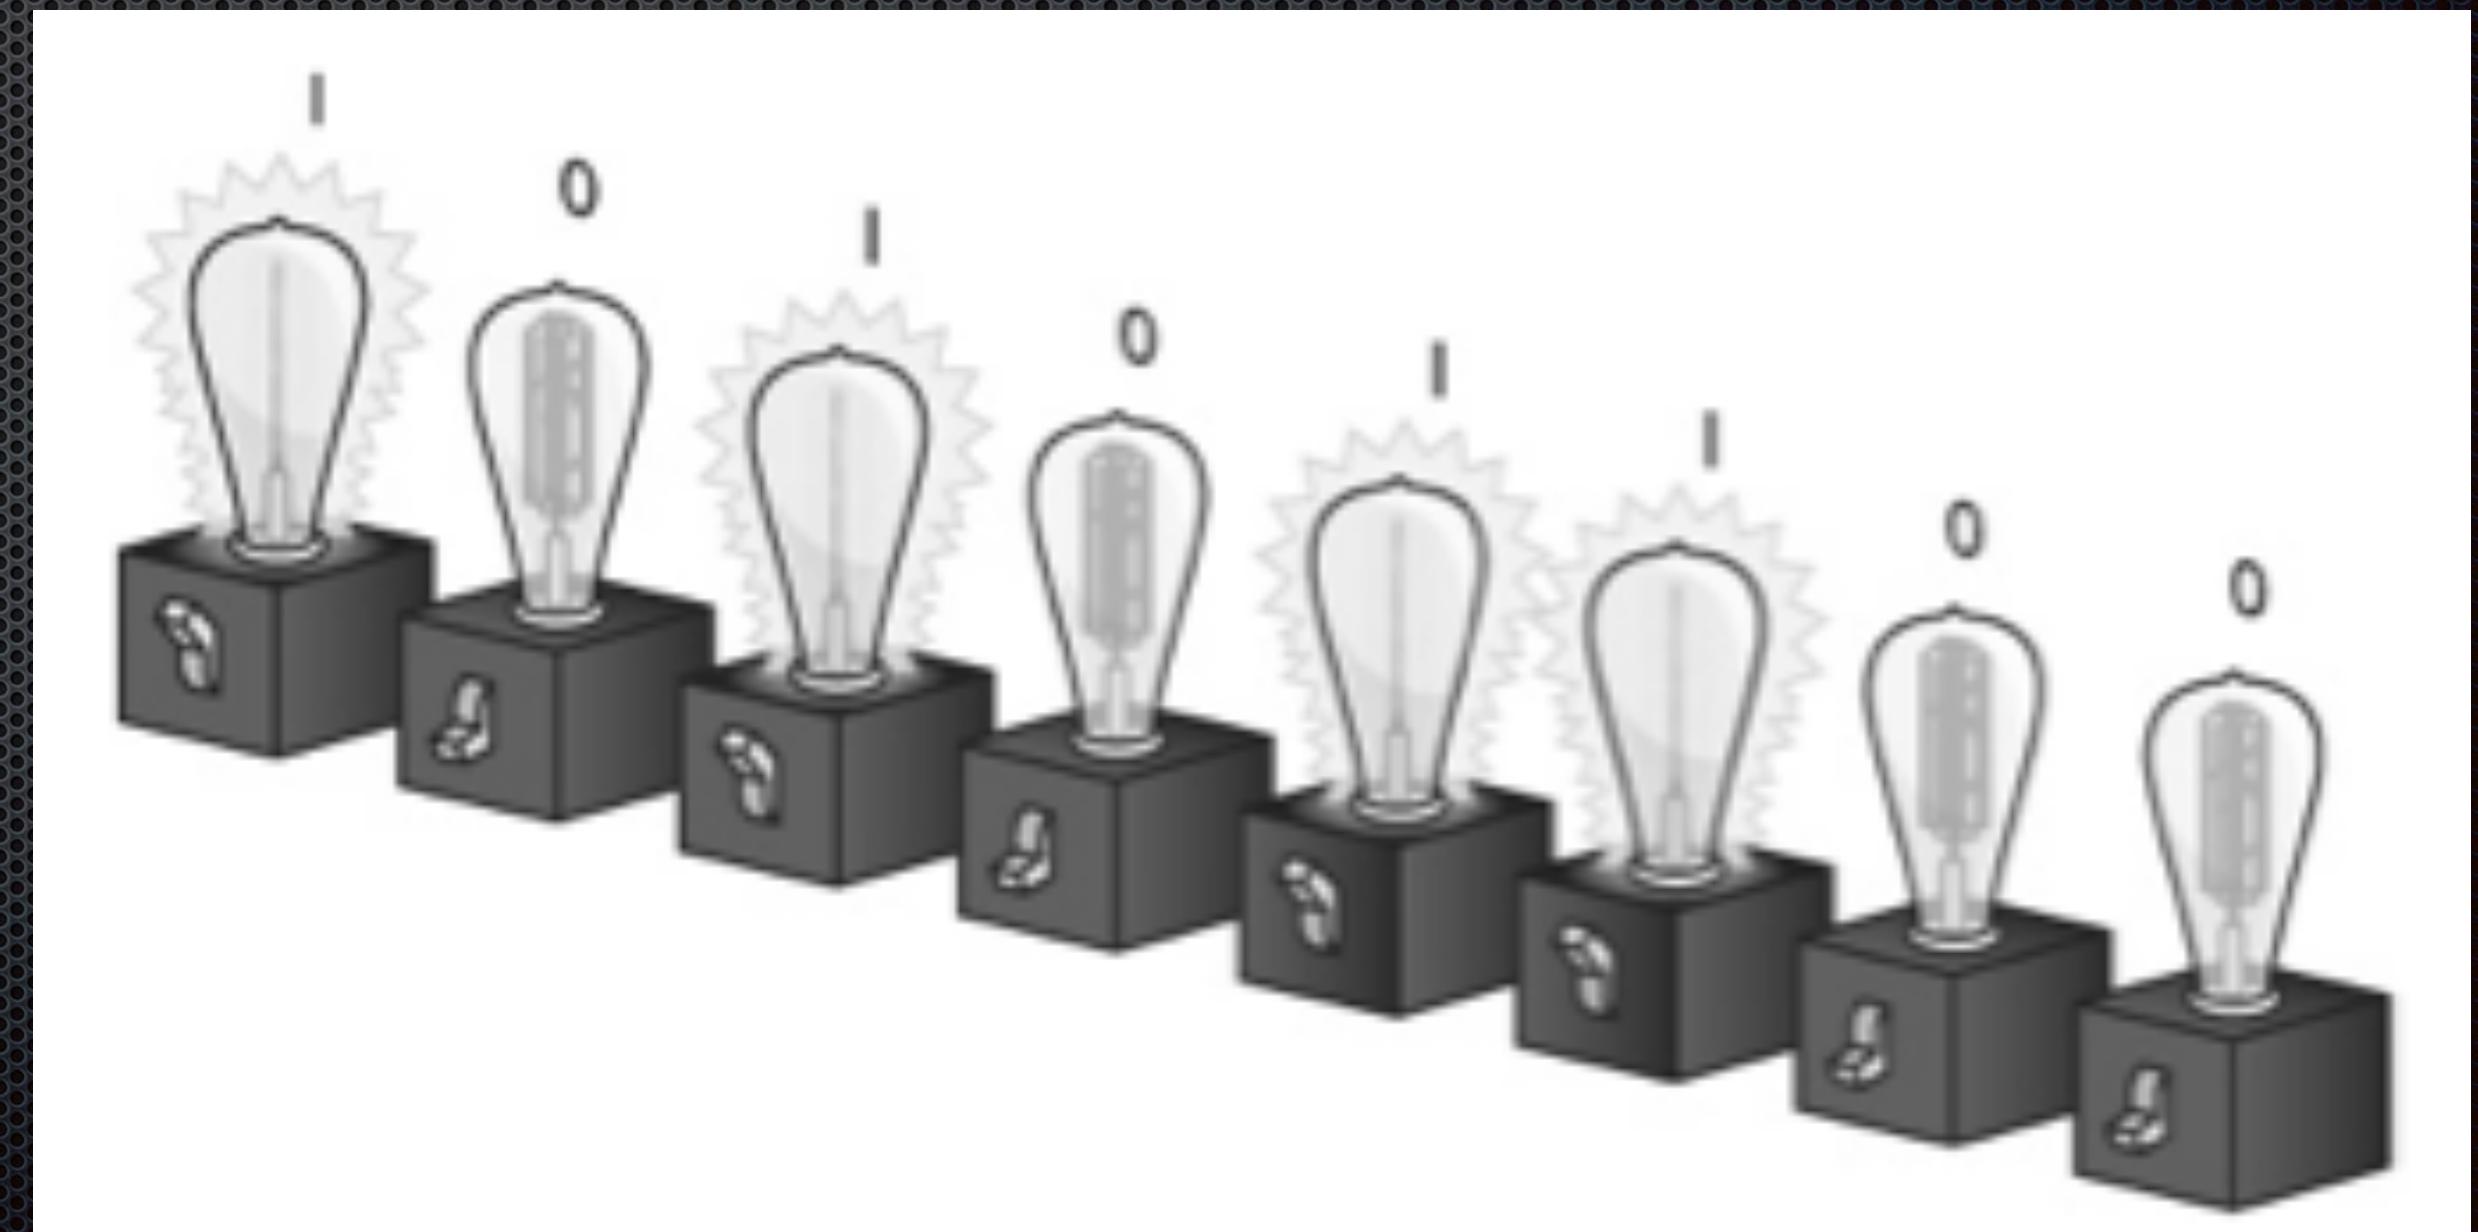


Data Bus

- Our light bulbs work as a data bus
- By turning bulbs on or off, different data sets can be represented
- Each bulb is visible to the guy in the box, and the people outside trying to talk to him
- Guy inside and people outside need to agree on what the patterns mean

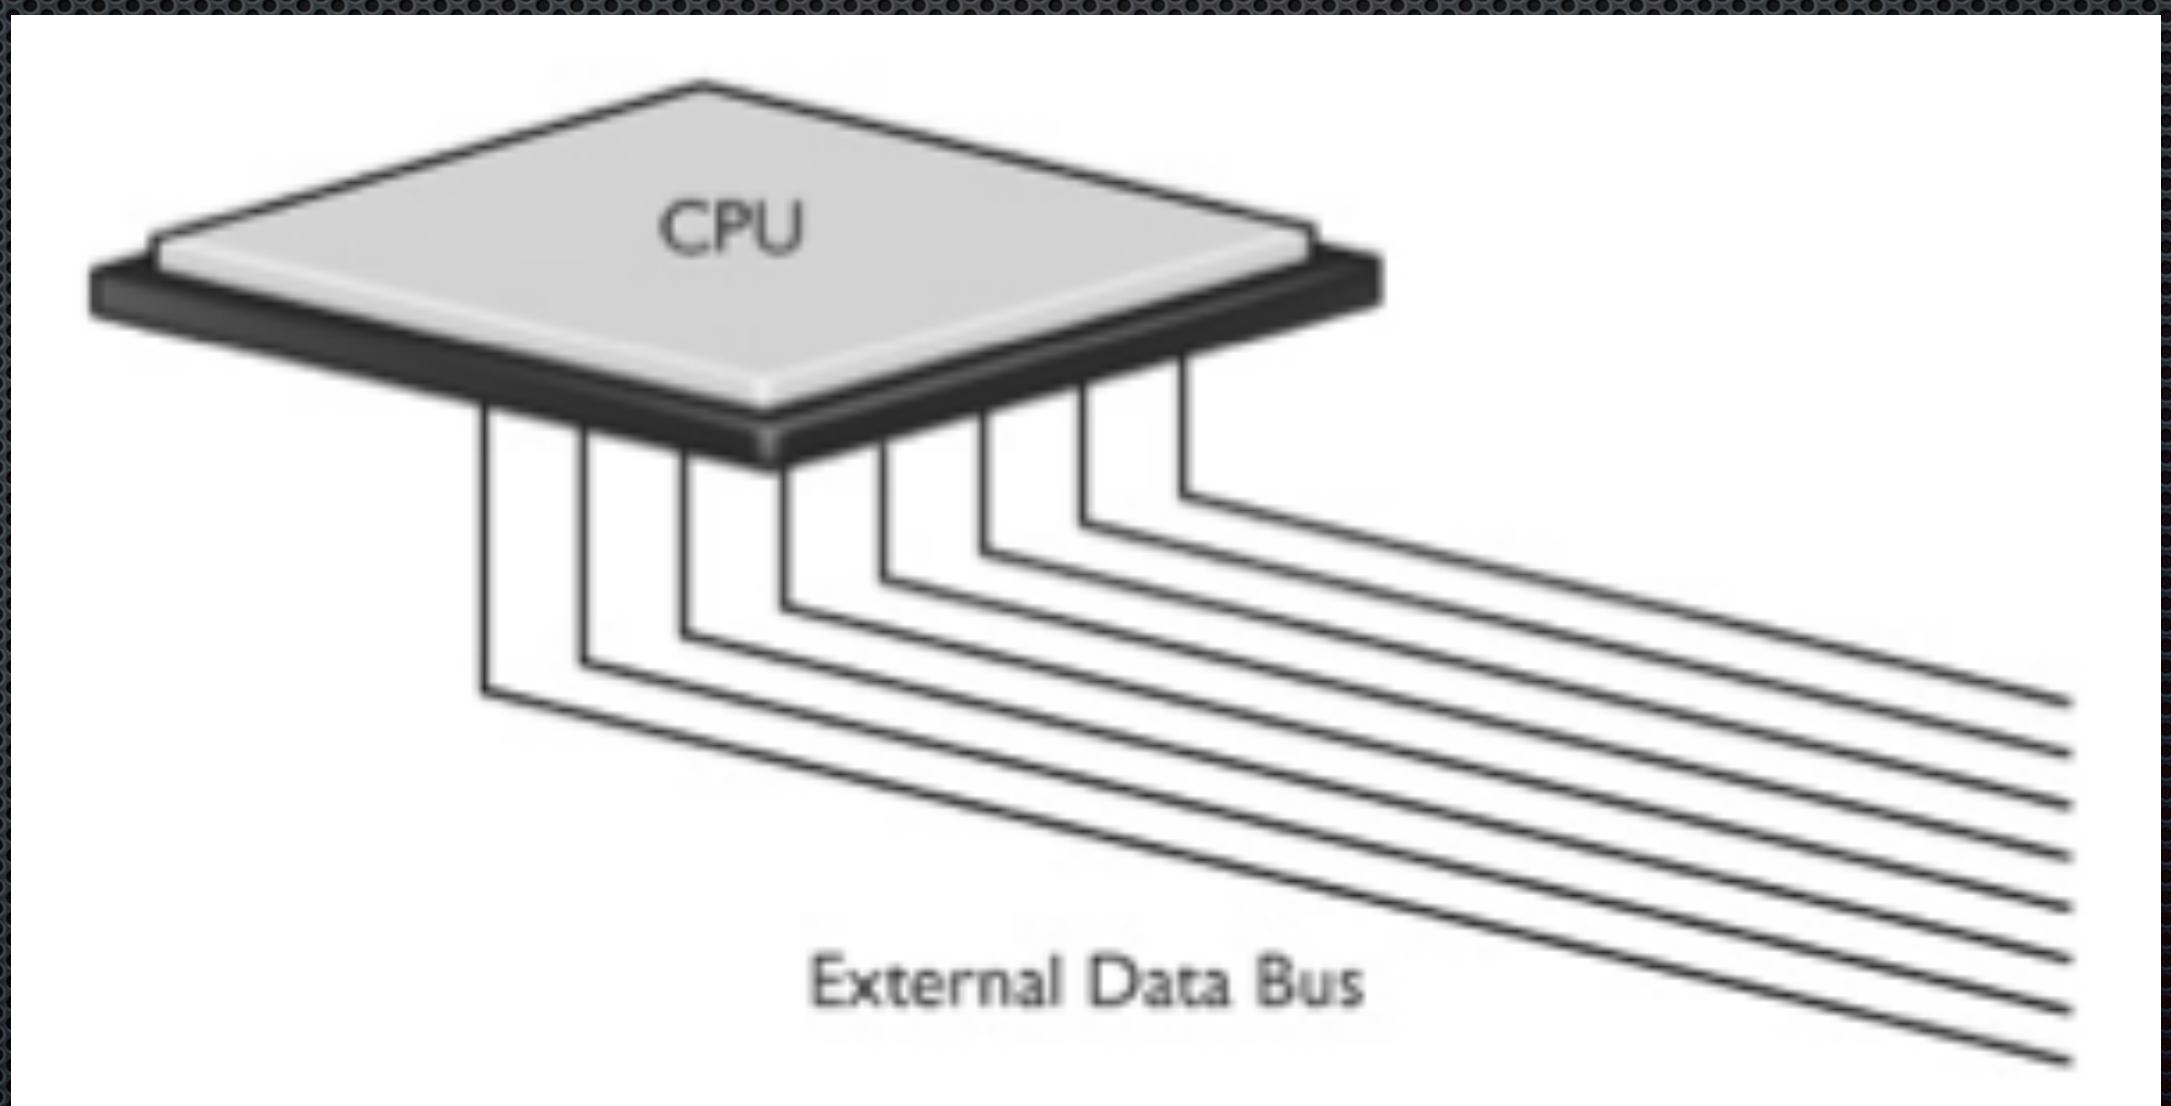
Binary Doesn't Go Away

- Back to the light bulbs, it's a pain to say on-off-off-on-etc.
- Binary will represent on and off
 - 1 = on
 - 0 = off



Address Bus

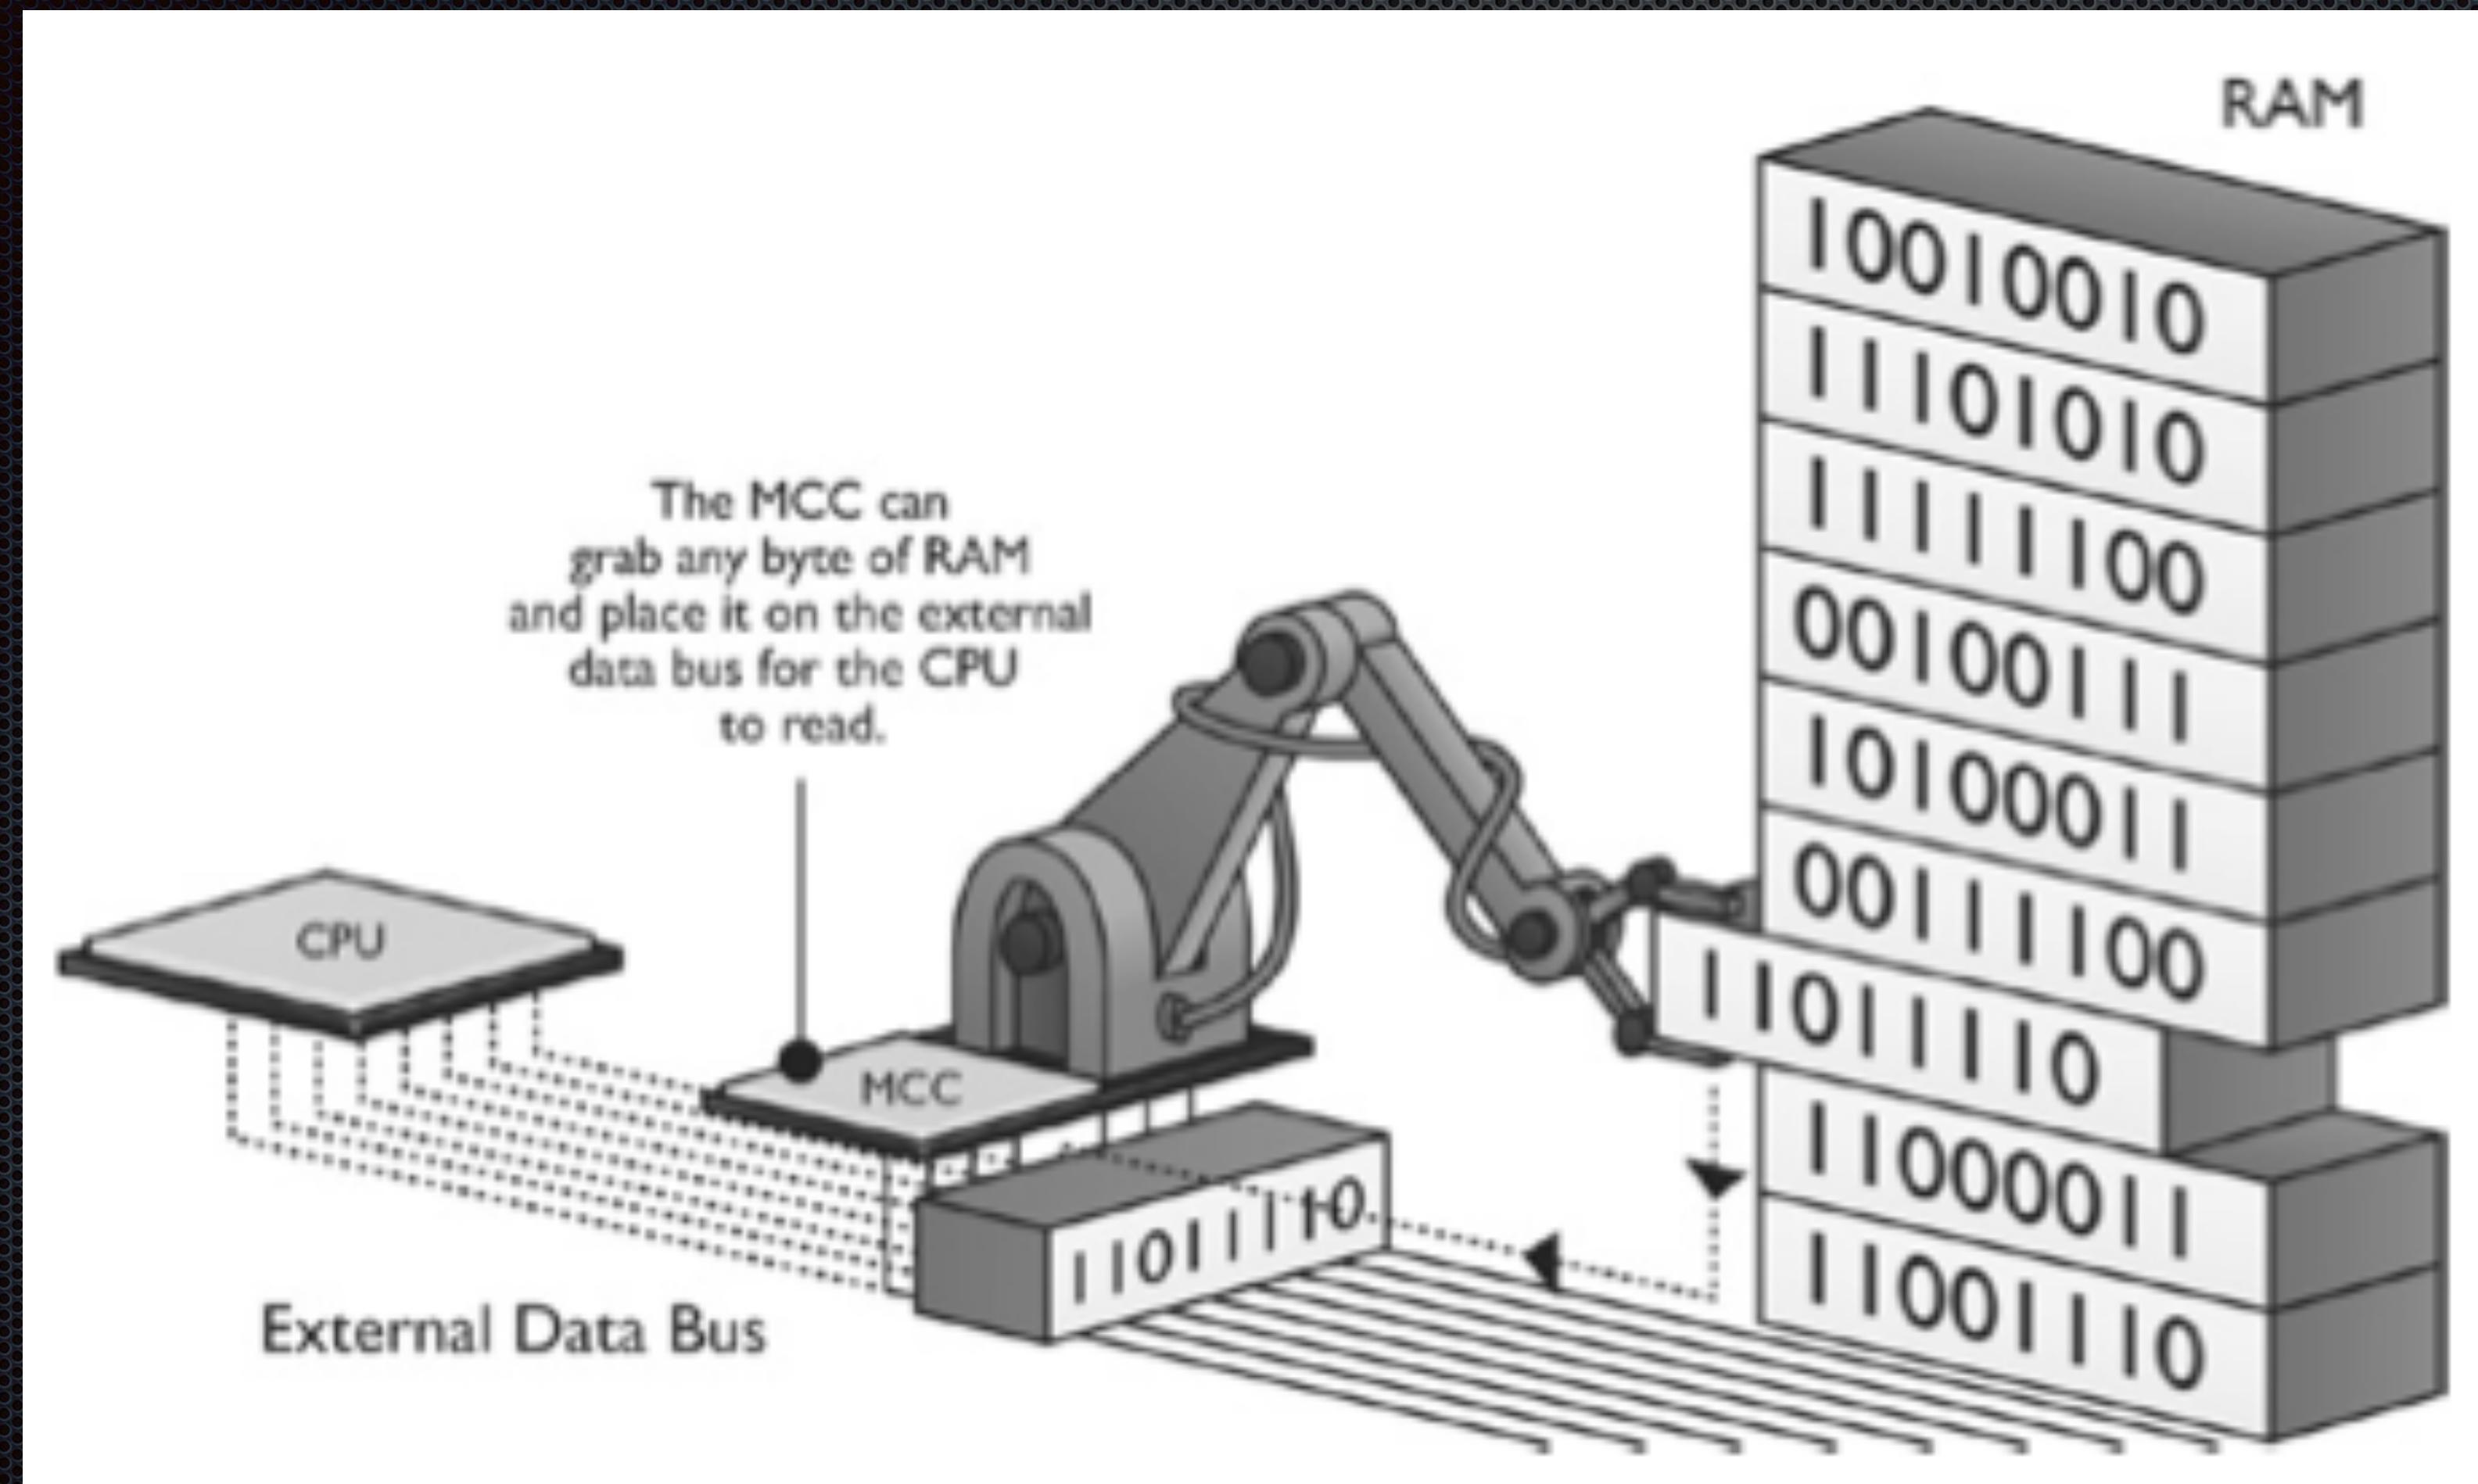
- So far, our computer is only a CPU and RAM
- Extend the EDB to support RAM



Address Bus

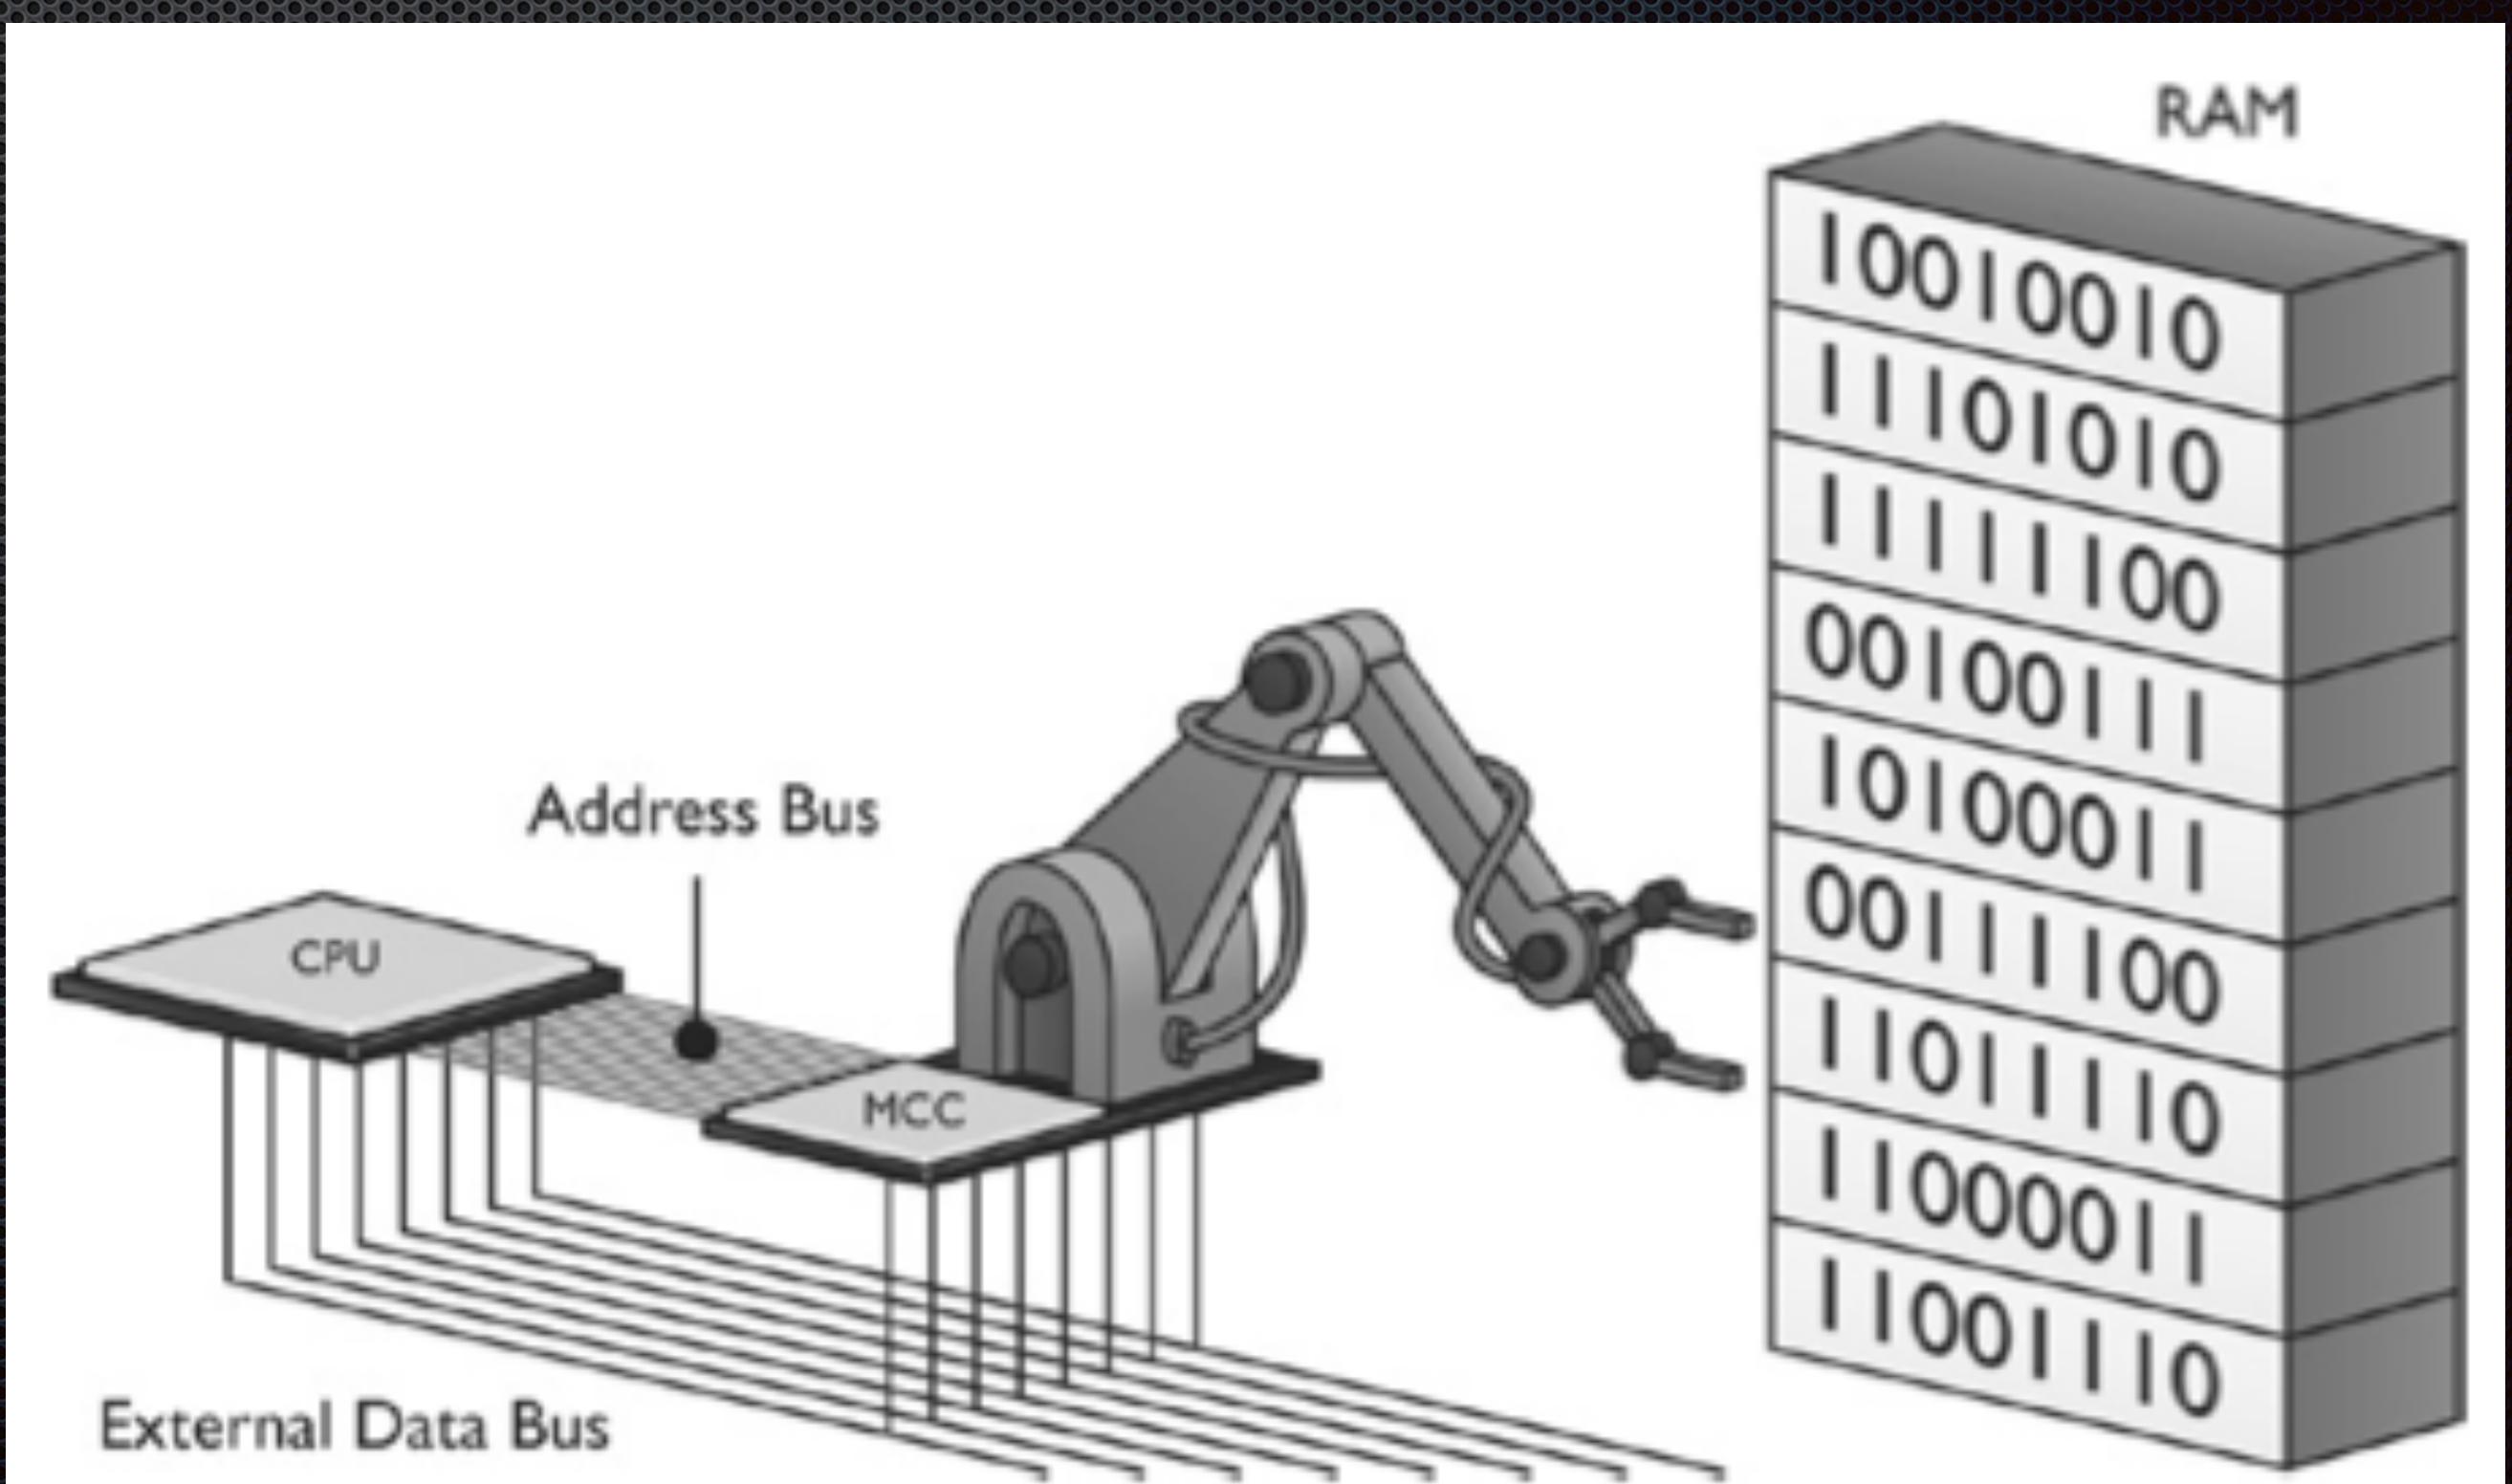
- It's not quite as simple as plugging RAM into the EDB
- There needs to be something in between for the processor to be able to select which row of RAM it needs
- Memory Controller Chip (MCC)

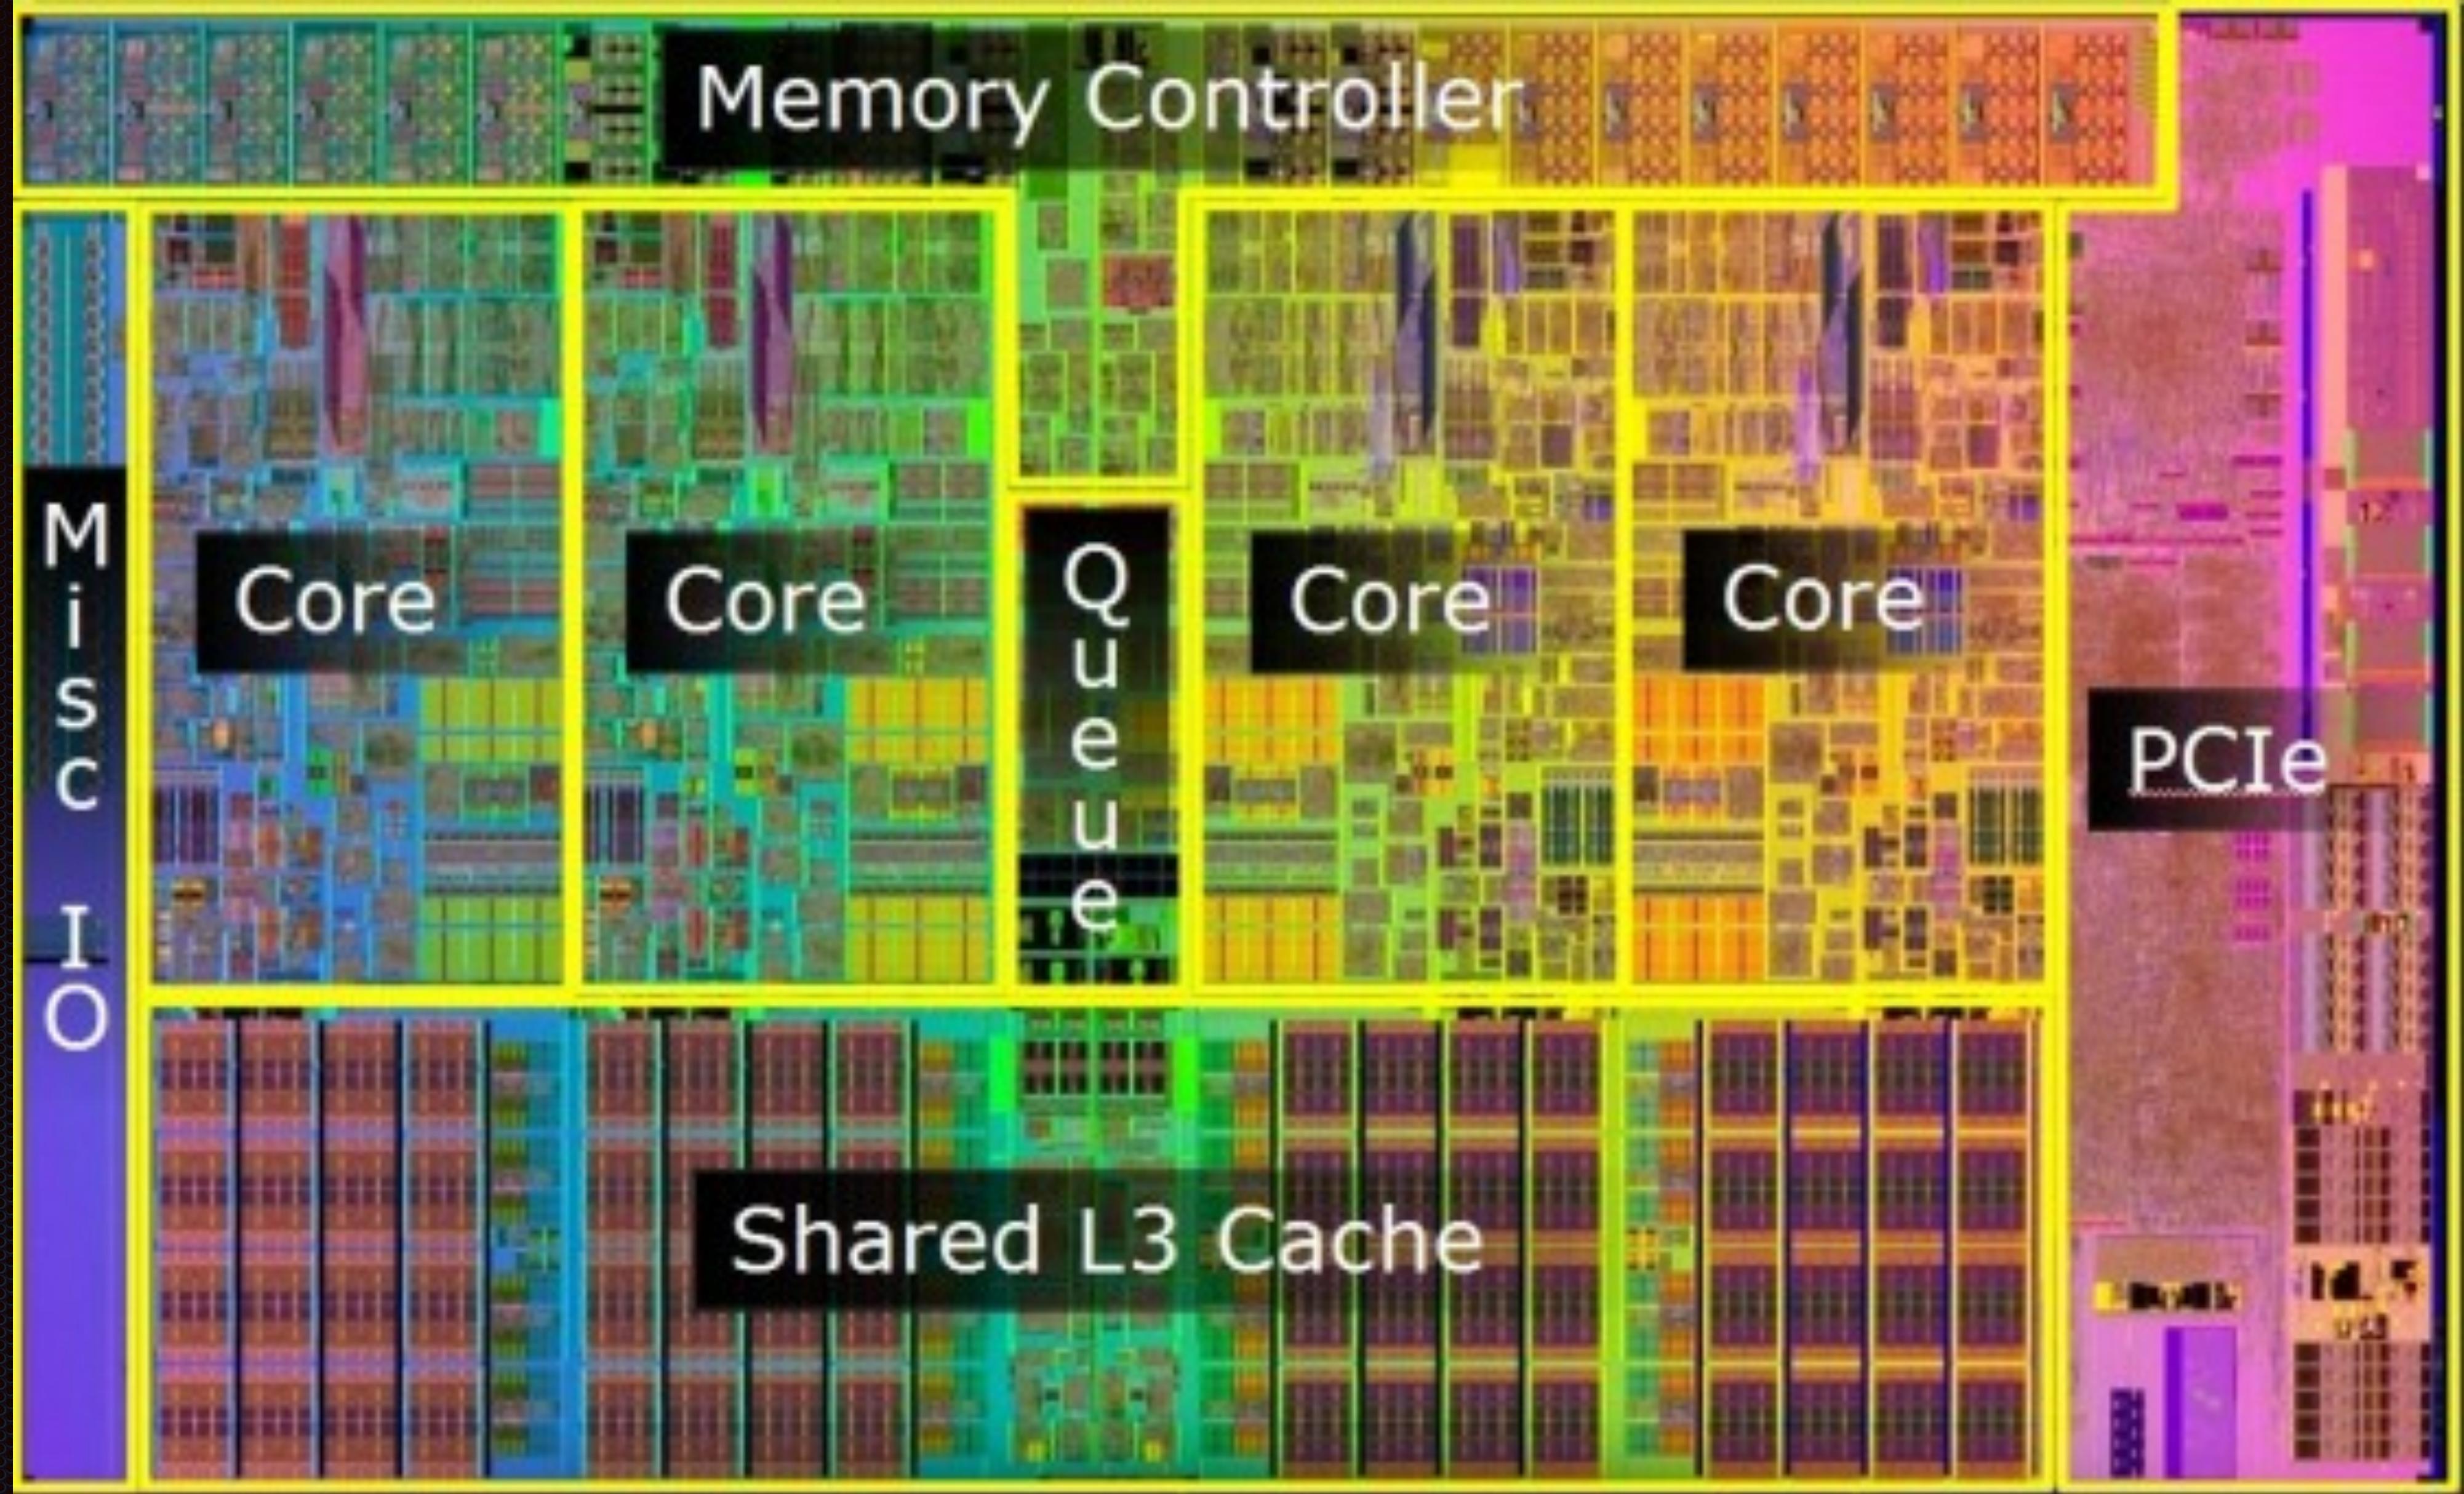
MCC



Address Bus

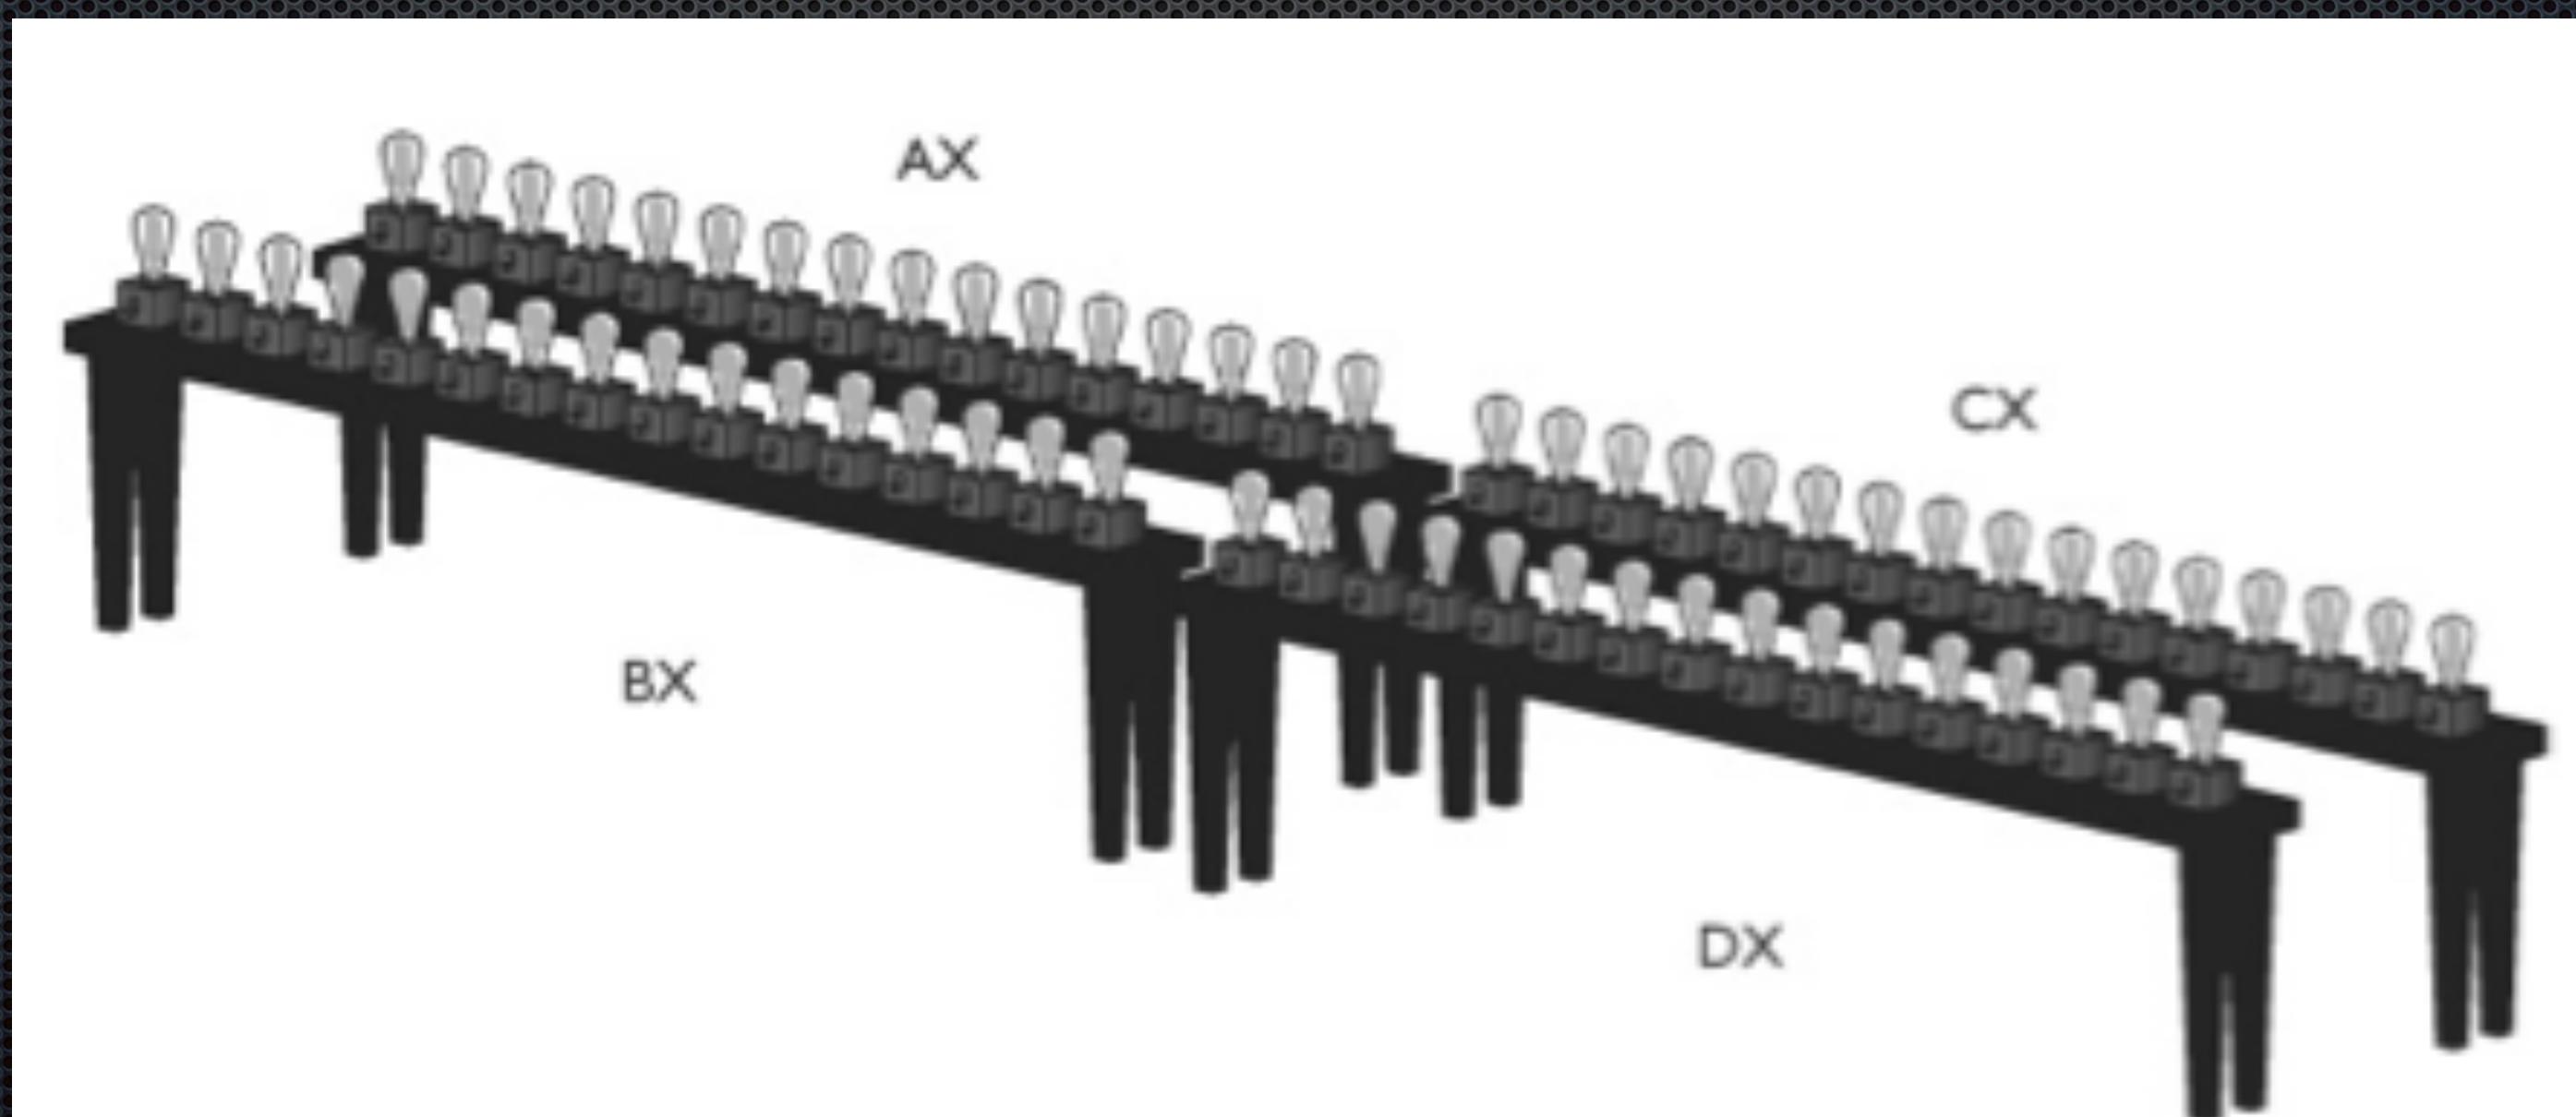
- MCC interfaces with EDB
- Address bus tells MCC which piece of RAM to grab





Going Inside the Box

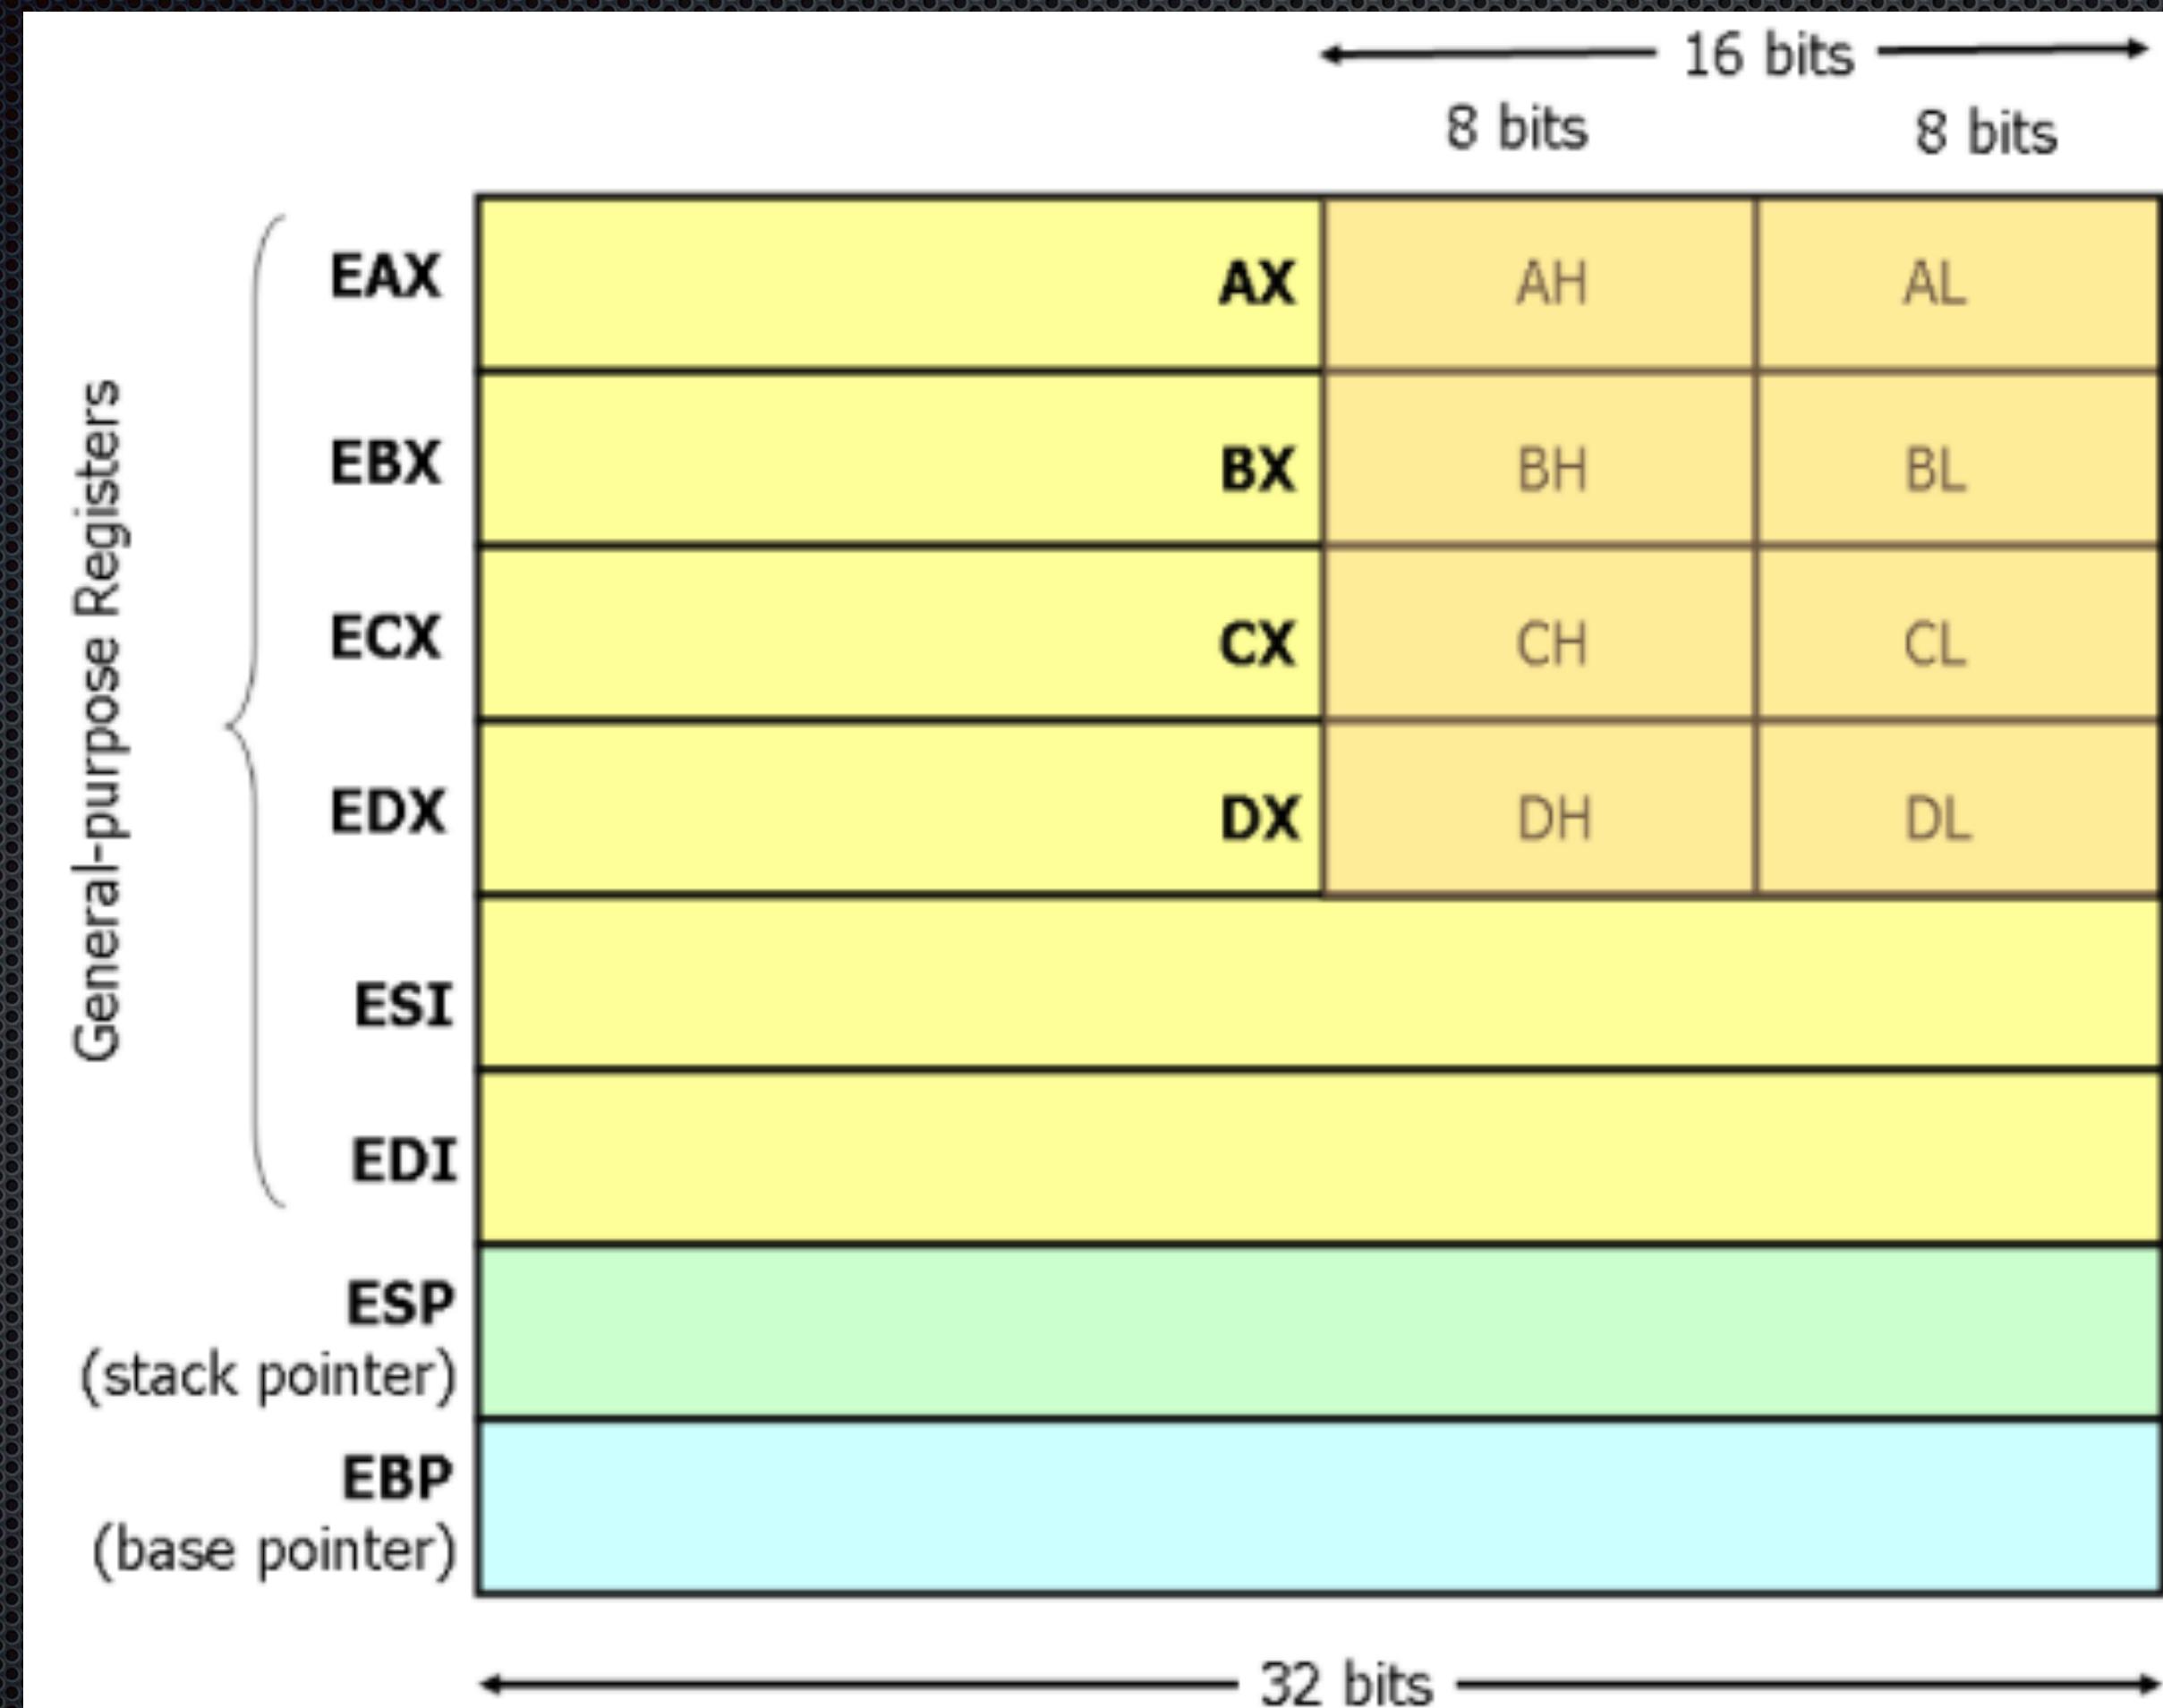
- The guy in the box needs work tables to handle the math and operations/instructions he's assigned
- These work tables are called registers



Registers

- All CPUs have them, some have different numbers of them
- 64-bit processors have more than registers than 32-bit processors
- They serve different functions, but let's focus on some common, general-purpose registers
- When RAM passes the processor an instruction or some information, it is stored on a workbench, or the register

32-Bit Registers



Remember our assembly?

```
0804842b <main>:
```

804842b:	8d 4c 24 04
804842f:	83 e4 f0
8048432:	ff 71 fc
8048435:	55

lea	0x4(%esp),%ecx
and	\$0xffffffff0,%esp
pushl	-0x4(%ecx)
push	%ebp

Another Quick Recap

- We compile C code into binary instructions for the processor to understand
- The program is loaded from the HDD to a location in memory
 - Each “row” of memory is addressed by a hexadecimal value (0xFFFFFDE73)
- The CPU retrieves instructions/data from memory and stores them in registers (workbenches)
 - The registers may hold variables, be used to do math/operations

Stack Buffer Overflow

- Open up the **stack_overflow_16.c** program and let's take a look at what it does
- The compilation instructions and execution instructions are in the code
- Run the program to see what it does, can you get an error?

```
15 #include <stdio.h>
16 #include <stdlib.h>
17
18 void echo(void)
19 {
20     printf("Enter some text:\n");
21     char buffer[16]; // 16 byte buffer stored on the stack
22     gets(buffer); //retrieve user input and store in the buffer
23     printf("%s\n", buffer); //print out the user input as a string
24     return;
25 }
26
27
28 void main(void)
29 {
30     echo();
31     exit(0);
32 }
```

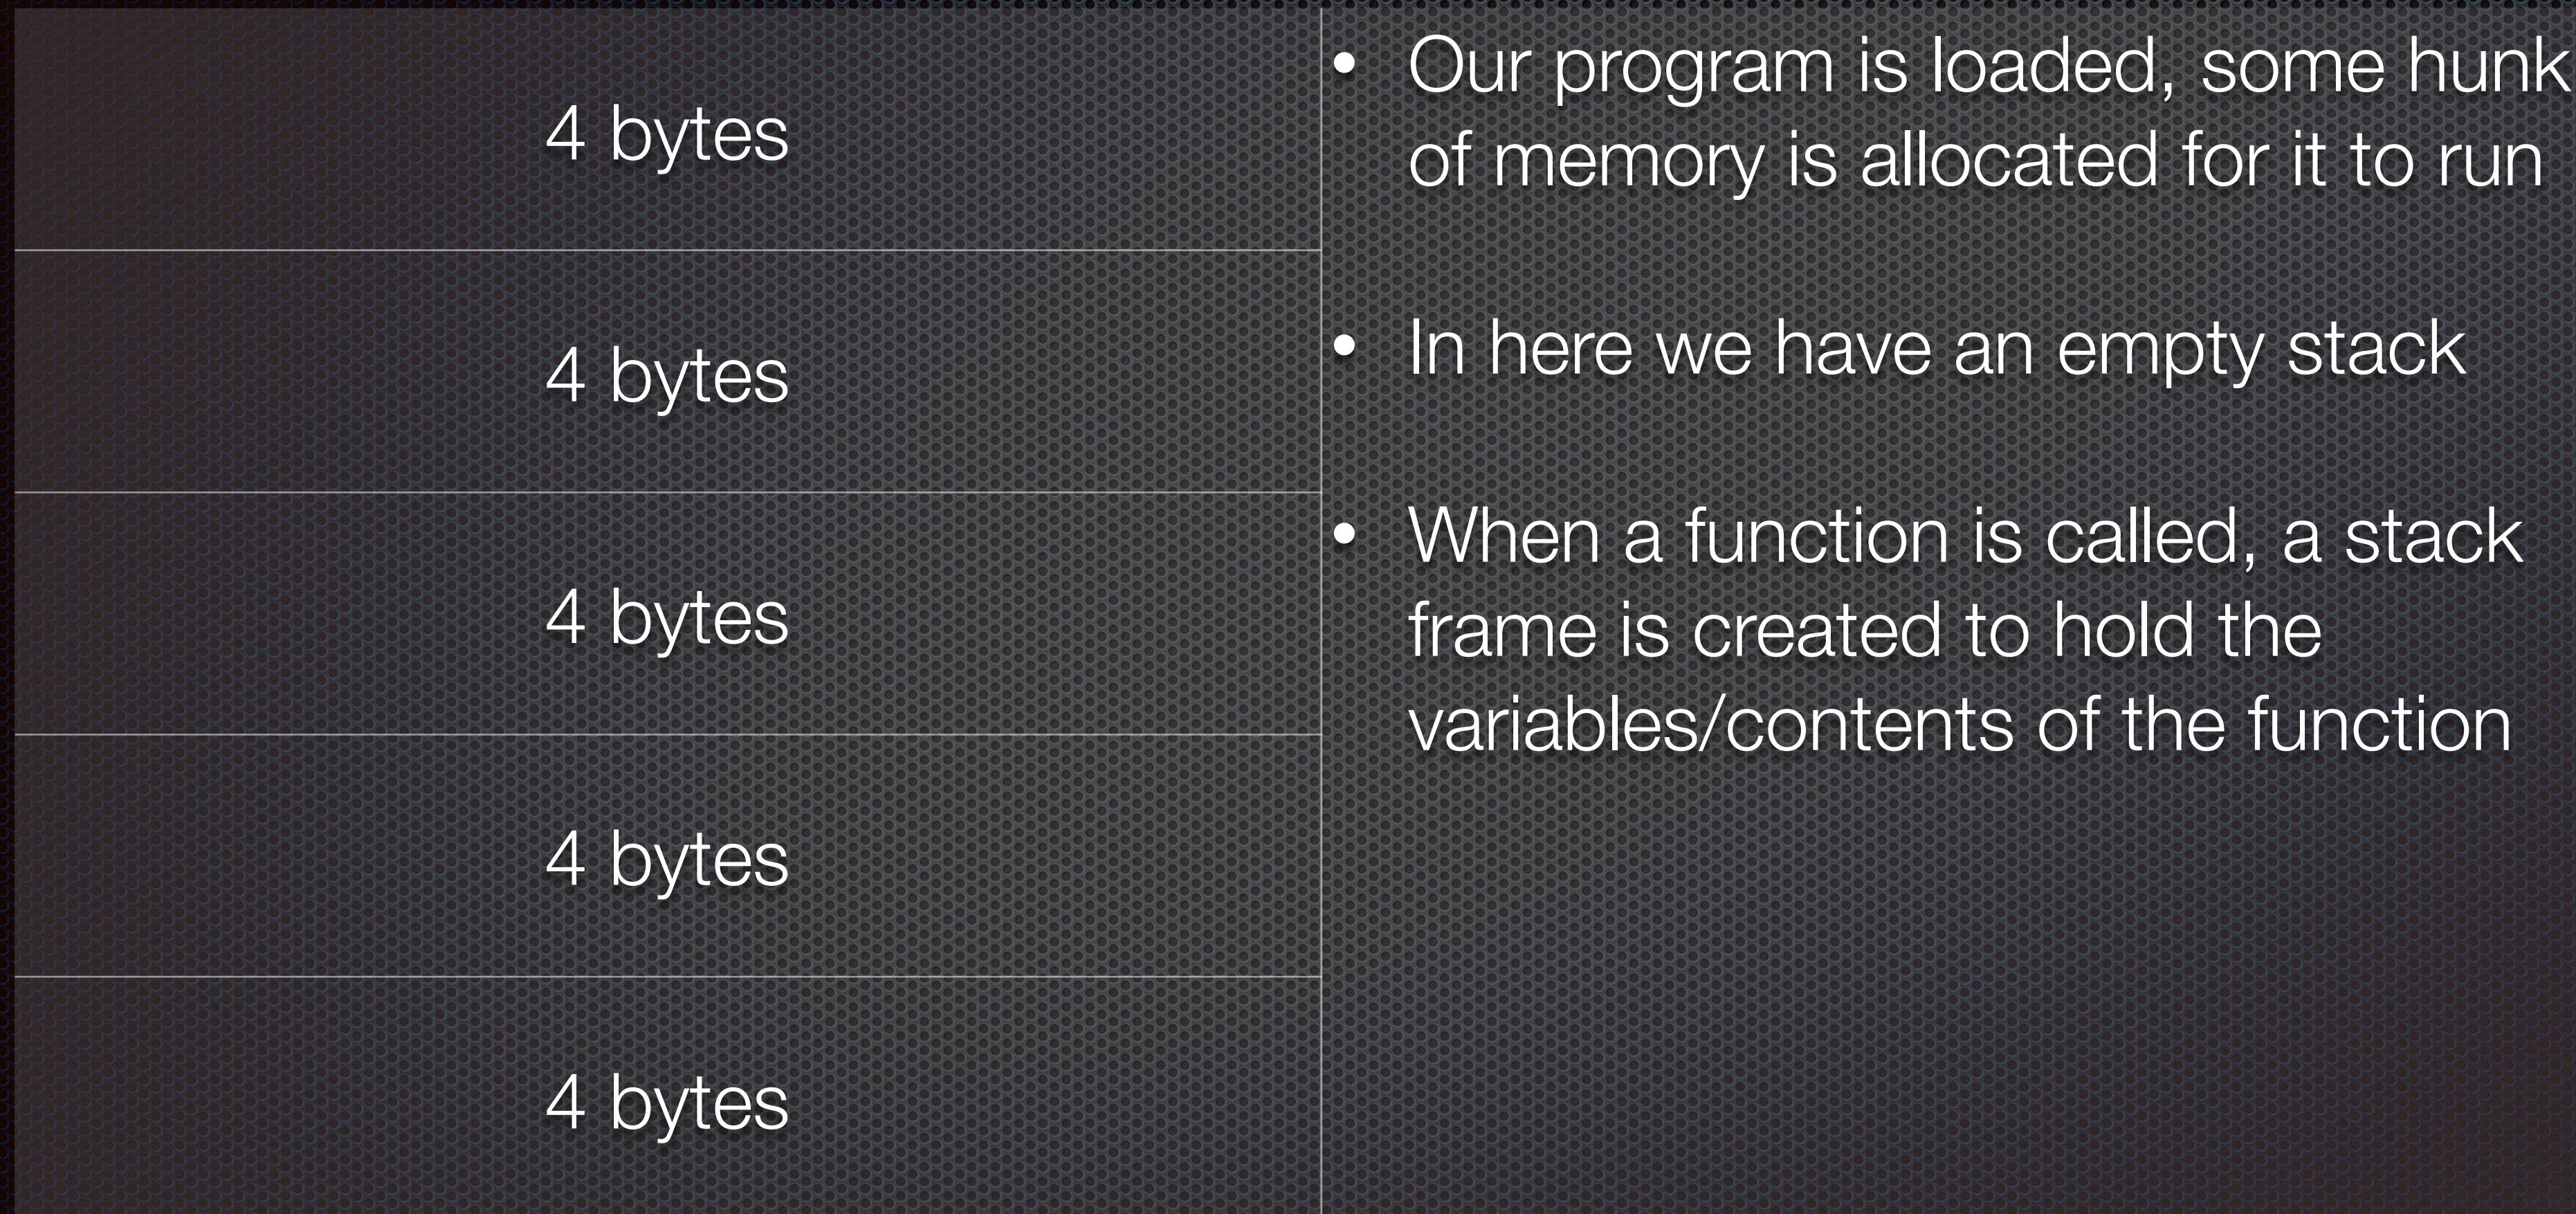
Manual Fuzzing (Testing Input)

- As you can see, if we enter a string larger than the buffer (16-bytes) we get a segmentation fault
- To see what is faulting, we need to open the program in a debugger:
 - `gdb ./stack_overflow_16.c`
 - `list 1,50`
 - `run`
- Cause the crash to happen again after running and look at the registers:
 - `info reg`

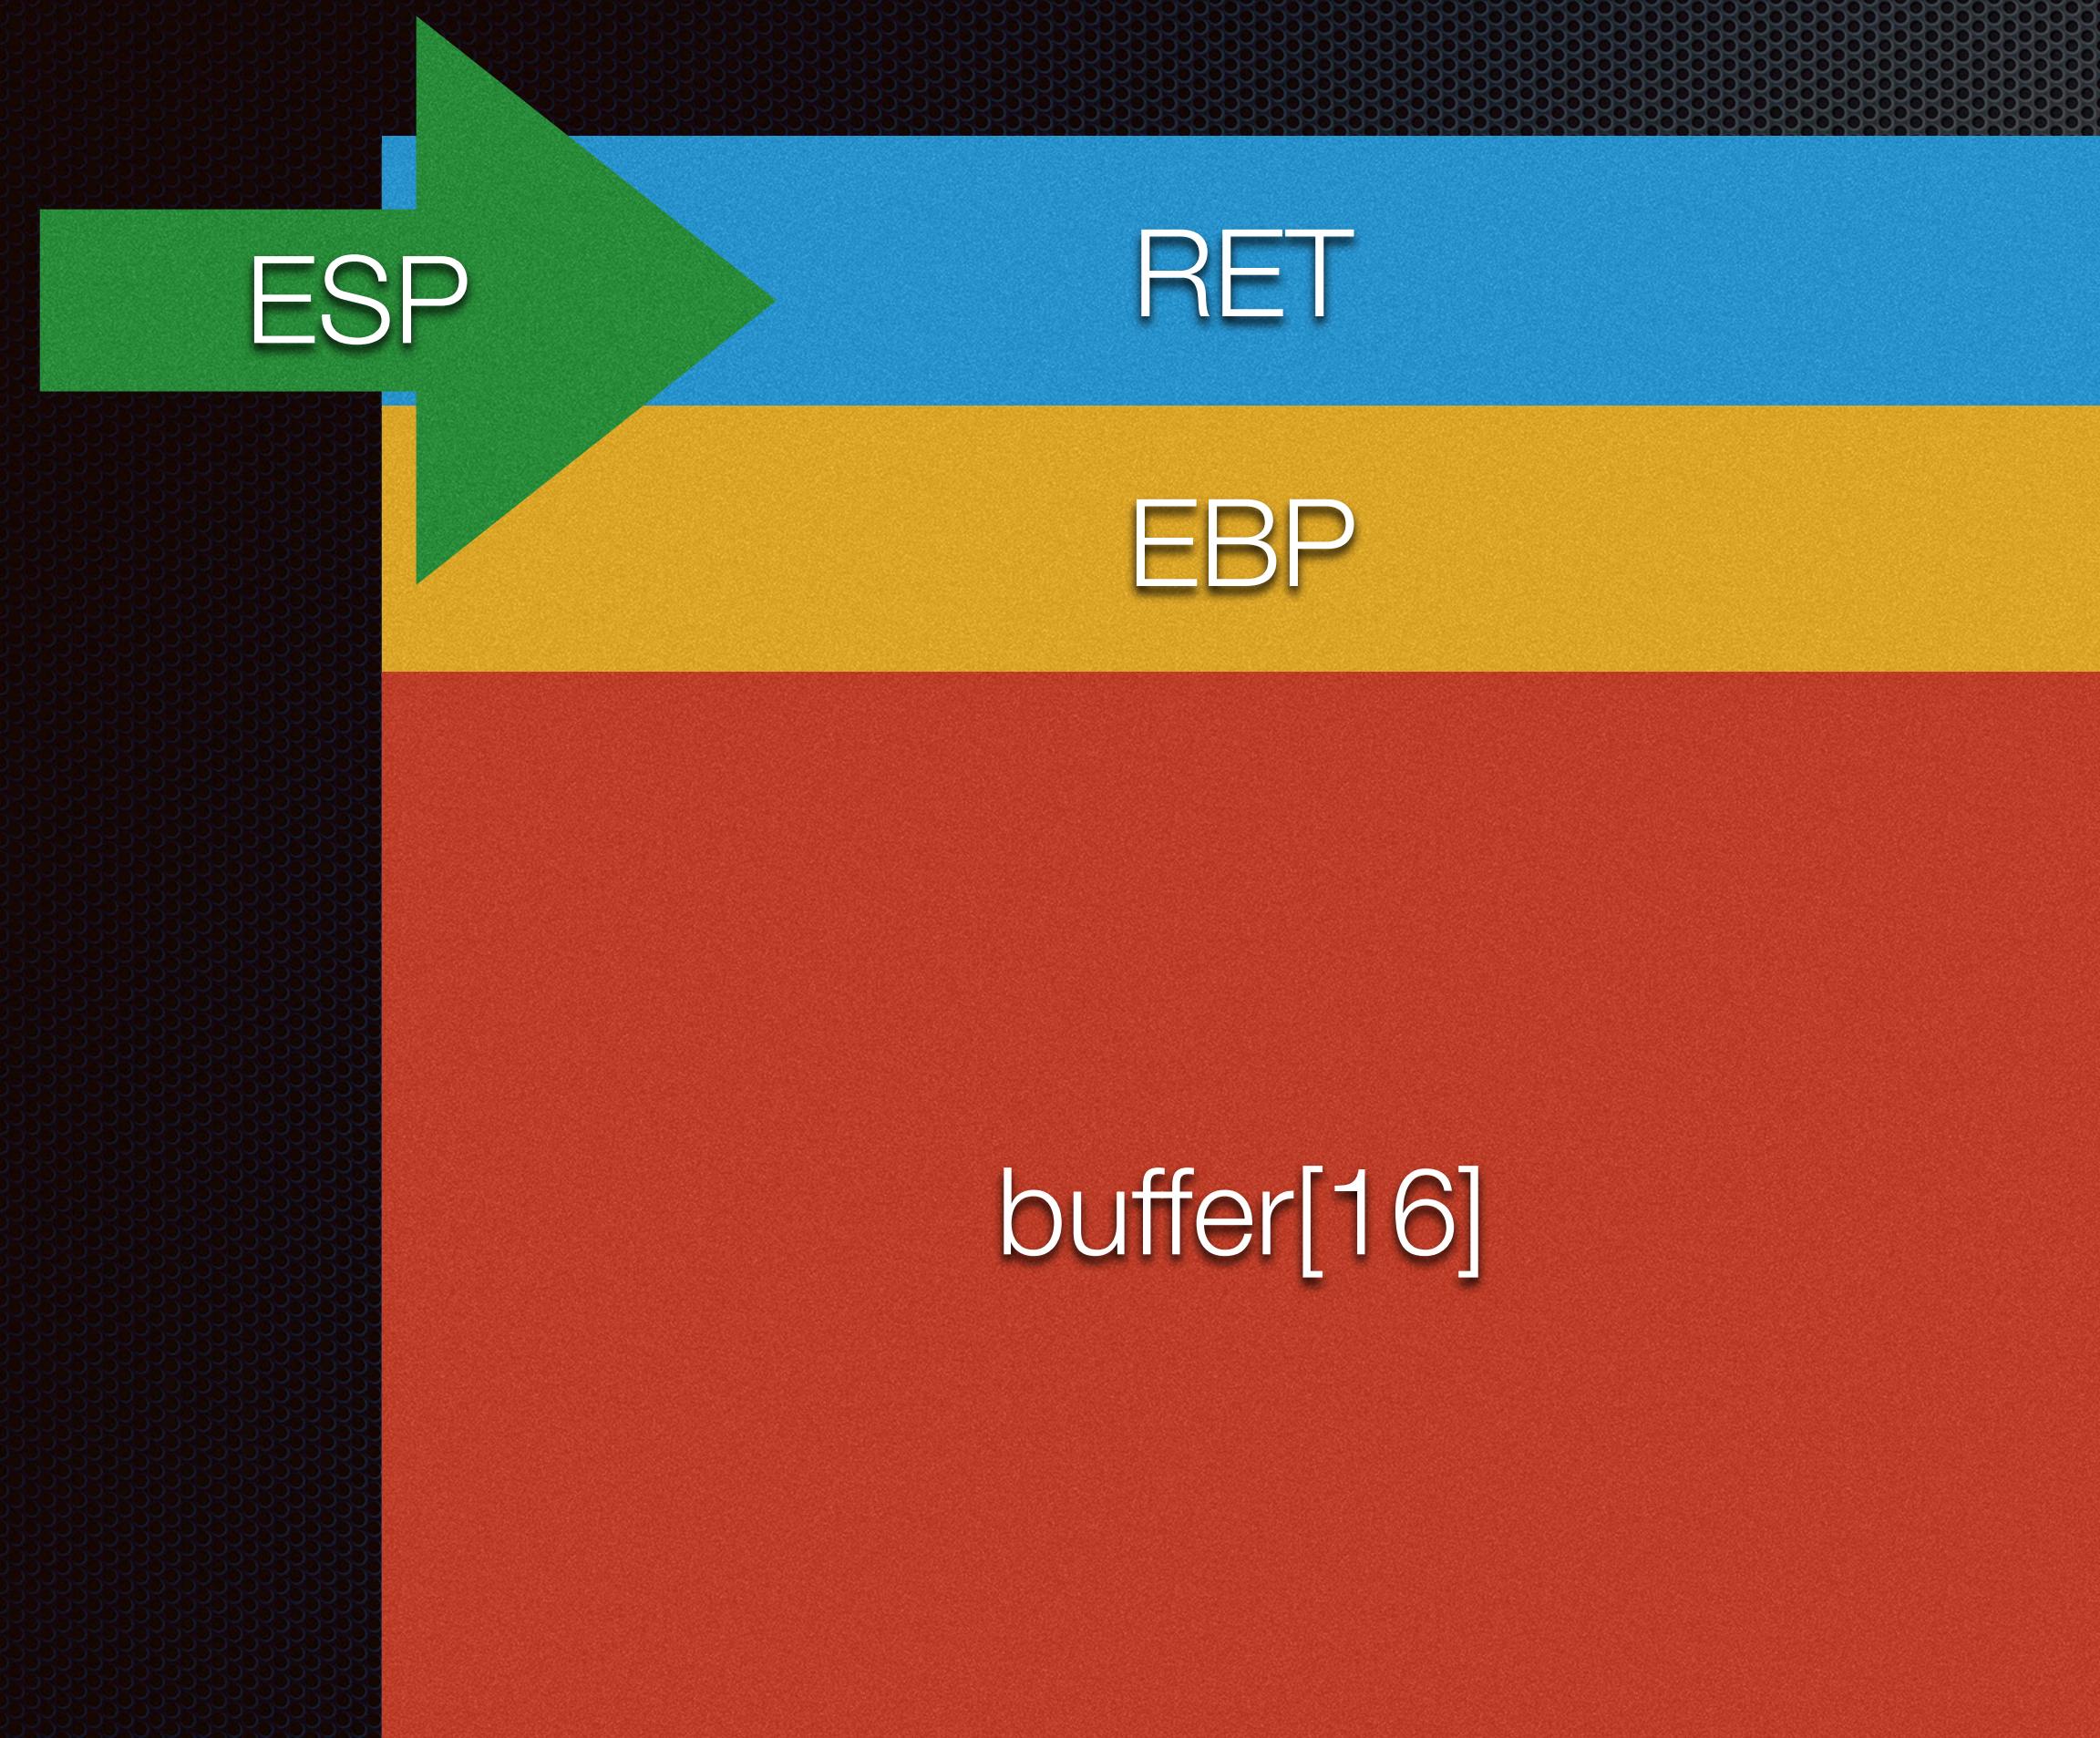
The Stack

- This is a section of your computers RAM where exploits can occur
- It is an arbitrary location in RAM that can grow/shrink depending on what your code needs
 - If you have a lot of variables in a function, the space in the stack will be larger
- When you load a program, one register (ESP) is used to control the stack

The Stack

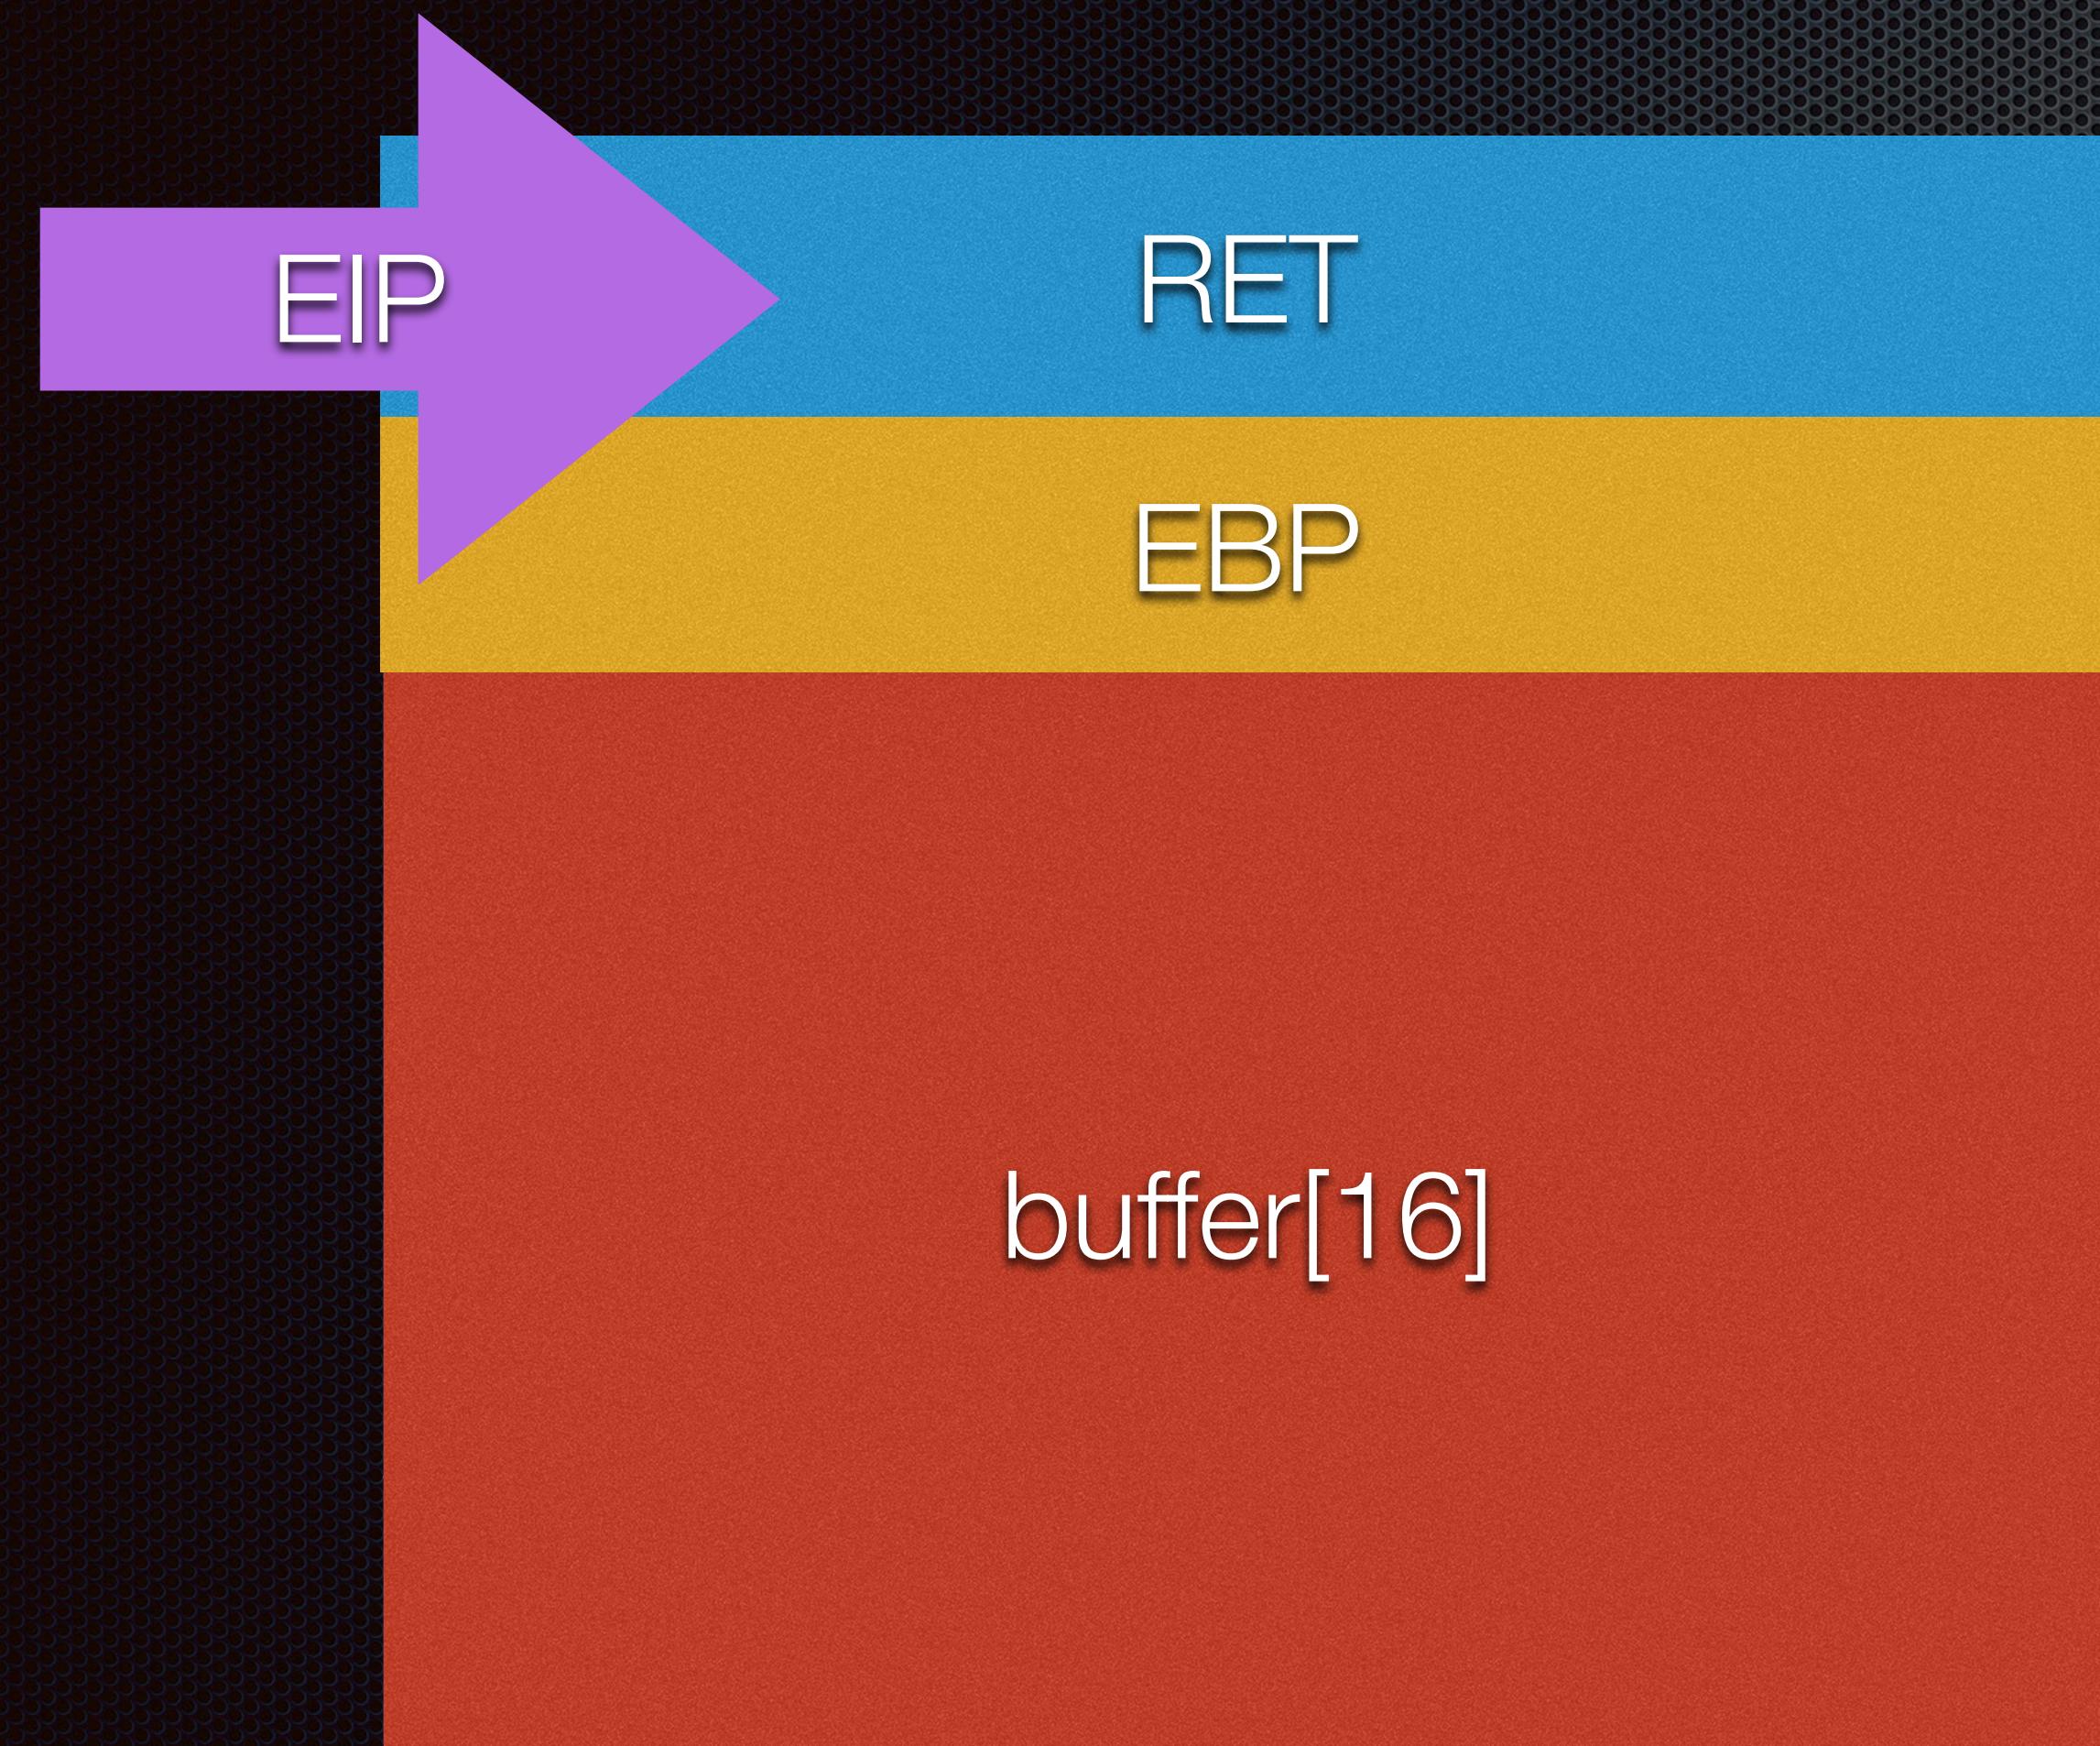


Calling a Function



- A register, ESP, keeps track of the location of the stack
- When the function is called, a return address is saved (RET)
- The base address of the stack is also saved (EBP)
- Then variables are added into the stack as well

Exiting a Function



- When the function exits, it's going to use a register, EIP
- EIP will get set to RET
- If the RET address is valid, the function exits

Stack Overflow

- Remember, EIP is the instruction pointer
 - Points to the address location in memory of the next instruction
 - In this case, it's trying to execute 0x48484848, an invalid address

```
Starting program: /root/Desktop/basic_exploitation/stack_overflow.out
Enter some text:
AAAABBBBCCCCDDDDEEEEFFFFGGGGHHHH
AAAABBBBCCCCDDDDEEEEFFFFGGGGHHHH

Program received signal SIGSEGV, Segmentation fault.
0x48484848 in ?? ()
(gdb) INFO REG
eax          0x21      33
ecx          0xffffffff -1
edx          0xf7fae878 -134551432
ebx          0xf7fad000 -134557696
esp          0xfffffd3f0 0xfffffd3f0
ebp          0x47474747 0x47474747
esi          0x0      0
edi          0x0      0
eip          0x48484848 0x48484848
```

What happened?

FFFF (4 bytes)

EEEE (4 bytes)

DDDD (4 bytes)

CCCC (4 bytes)

BBBB (4 bytes)

AAAA (4 bytes)

- Our one variable was only supposed to hold 16 bytes
- We wrote much more than that and ended up eventually overwriting the RET address
- EIP tries to execute FFFF and fails when the function exits

The Real Crash Happened at HHHH

- What's the deal, we crashed at HHHH, your slide showed FFFF?!?
- The way memory is managed and the compiler works, it's not always 100% picture perfect
- There are optimizations that occur and memory is organized in what should be the most efficient manner
- Sometimes segments of memory get aligned differently
- Exploitation is sometimes an art, not a science :)

Take a moment...we have manipulated what is in EIP
We choose what the program does next :)
This is a proof-of-concept (POC) exploit at this point!

What's on the stack?

- We'll look at the hex words leading up to ESP, the stack pointer
- The highlighted text is the start of our input (0xffffd3d0)

```
(gdb) x /64wx $esp-32
0xfffffd3d0: 0x41414141 0x42424242 0x43434343 0x44444444
0xfffffd3e0: 0x45454545 0x46464646 0x47474747 0x48484848
0xfffffd3f0: 0xf7fad300 0xfffffd410 0x00000000 0xf7e1da63
0xfffffd400: 0x080484c0 0x00000000 0x00000000 0xf7e1da63
0xfffffd410: 0x00000001 0xfffffd4a4 0xfffffd4ac 0xf7feb79a
0xfffffd420: 0x00000001 0xfffffd4a4 0xfffffd444 0x08049780
0xfffffd430: 0x0804821c 0xf7fad000 0x00000000 0x00000000
0xfffffd440: 0x00000000 0xdbe9f272 0xe7f59662 0x00000000
0xfffffd450: 0x00000000 0x00000000 0x00000001 0x08048360
0xfffffd460: 0x00000000 0xf7ff0fe0 0xf7e1d979 0xf7ffd000
0xfffffd470: 0x00000001 0x08048360 0x00000000 0x08048381
0xfffffd480: 0x08048492 0x00000001 0xfffffd4a4 0x080484c0
0xfffffd490: 0x08048530 0xf7febcb0 0xfffffd49c 0x0000001c
0xfffffd4a0: 0x00000001 0xfffffd60c 0x00000000 0xfffffd640
0xfffffd4b0: 0xfffffd64b 0xfffffd65c 0xfffffd66f 0xfffffd69a
0xfffffd4c0: 0xfffffd6ab 0xfffffd6c2 0xfffffd6d2 0xfffffd6dd
(gdb) █
```

Making a Useful Payload

- Right now, our “payload” is a series of letters (AAAAABBBBB[...]HHHH)
- We can replace those letters in our payload with instructions to perform some other action on the system
 - HHHH is the instruction that should be executed next, but we could change that to the start of our payload (0xffffd3d0) so the payload gets executed
- Take a look at **msfvenom** a bit later
 - Built-in tool to help us generate a payload to be used

Let's Modify the Program a Bit

- 16 bytes isn't a whole lot of space to have to exploit or introduce code
- We're going to switch to **stack_overflow_256.c**
- Enter in a large number of characters again to see if we can get a crash and control EIP again
- perl -e 'print "A"x256,"BBBBCCCCDDDDEEEEFFFFGGGGHHHH"'

```
18 void echo(void)
19 {
20     printf("Enter some text:\n");
21     char buffer[256]; // 256 byte buffer stored on the stack
22     gets(buffer); //retrieve user input and store in the buffer
23     printf("%s\n", buffer); //print out the user input as a string
24     return;
25 }
```

Fuzzing Again

- We did 256 A (256 bytes)
- BBBB (4 bytes)
- CCCC (4 bytes)
- DDDD (4 bytes)
- EEEE (4 bytes) <- what overwrote EIP
- Therefore we have 268 bytes of shellcode ($256+4+4+4$), 4 bytes for EIP
- Total payload length = 272 bytes

Shellcode

- Shellcode is code represented in hex that can be interpreted by the CPU
 - This is where we get to tell the system what to do
- Here is some simple shellcode to print out “y0u sp34k 1337 ?” (50 bytes):
 - eb 19 31 c0 b0 04 31 db b3 01 59 31 d2 b2 12 cd 80 31 c0 b0 01 31 db b3 01 cd 80 e8 e2 ff ff ff 20 79 30 75 20 73 70 33 34 6b 20 31 33 33 37 20 3f 20
- We'll look at generating our own at some point

New Shellcode

- The buffer is 256 bytes that can hold real data
- The shellcode to print “y0u sp34k 1337 ?” is 50 bytes
- So, 256-50 leaves us with 206 bytes of free space
 - How do we fill up 206 bytes of free space?
- Add on 12 bytes of nothing and 4 bytes for EIP
- We should have 272 bytes in the end

NOP Sled

- When you have free space to fill up in some shellcode, you can use a NOP sled
- It's just a series of 1 byte instructions that basically tell the program to continue on
- On a normal 32-bit processor, the instruction for NOP is 0x90
- `perl -e ‘print “\x90”x206’` will give you 90 bytes of NOP

MSFVENOM

- List the available payloads:
 - msfvenom --help
 - msfvenom --list payloads
- Choose a payload of your liking and we can use it:
 - msfvenom -p linux/x86/exec CMD=ls -f c -b '\x00'

Redirection

- Sometimes, we don't need to execute our own shellcode
- We might just want the program to do something it is intended to
 - Give us administrative privileges
 - Say our serial number/product key is genuine
 - Move money from one bank account to another
- If the functionality is already built in, let's just use it!

stack_overflow_redirection.c

- Let's start out like we've done before
- Run the program, see what it does
- Can you get it to crash?
- What functions are being called?

```
stack_overflow_redirection.c *
```

```
1  /*
2   * Compile instructions (x32):
3   * gcc -g -fno-stack-protector -z execstack stack_overflow_redirection.c -o stack_overflow_redirection
4   *
5   * Compile instructions (x64):
6   * gcc -g -m32 -fno-stack-protector -z execstack stack_overflow_redirection.c -o stack_overflow_redirection
7   *
8   * Run instructions:
9   * ./stack_overflow_redirection.out
10  *
11  * Debug instructions:
12  * gdb ./stack_overflow_redirection.out
13 */
14
15 #include <stdio.h>
16 #include <stdlib.h>
17 #include <string.h>
18
19 void lockOutUser(void)
20 {
21     // If the user enters in the wrong password, they must be a hacker. Lock the account.
22     printf("You entered the wrong password. Account locked out!\n");
23 }
24
25 void giveUserRoot(void)
26 {
27     // If we're here, the user must be legit. Full admin privileges.
28     printf("Root access granted!\n");
29 }
30
31 int passCheck(void)
32 {
33     char buff[16]; // Passwords can be up to 15 characters
34     int validPass = 0;
35
36     // Prompt the user for a password and store it in "buff"
37     printf("Enter the password : \n");
38     gets(buff);
39
40     // Password checking routine
41     if(strcmp(buff, "gen_cyber=1337"))
42     {
43         // Incorrect
44         printf ("Wrong Password. \n");
45         validPass = 0;
46     }
47     else
48     {
49         // Correct
50         printf ("Correct Password! \n");
51         validPass = 1;
52     }
53
54     return validPass;
55 }
56
57 int main (void)
58 {
59     if (passCheck() == 1)
60         giveUserRoot();
61     else
62         lockOutUser();
63
64     exit(0);
65 }
```

So you can control EIP, huh?

- Bringing this full circle
- The type of exploit (stack buffer overflow) isn't really different than the last
- Rather than redirecting input to our shellcode that we chose, this time we are going to redirect execution (EIP) to another spot in the real program
 - We have a really small buffer this time!
- Starting to build out your swiss army knife of exploitation
- All we need to do next is to find the memory address of a valid function...but how?

Helpful GDB Commands

- List all of the functions within a program
 - `info func`
- Set a breakpoint (stop the program) when a function is called
 - `break <func name>`
- List breakpoints
 - `info break`
- Disassemble Function
 - `disas <func name>`

Let's take a look...

- Generate another long fuzzing string (maybe twice the length of the buffer)
- Find out where it is crashing like we've done in the past
- Once you know how many bytes you need to overwrite EIP, look for a call to the desired function
 - Replace EIP with the call address

```
(gdb) disass main
Dump of assembler code for function main:
0x08048534 <+0>: lea    0x4(%esp),%ecx
0x08048538 <+4>: and    $0xffffffff0,%esp
0x0804853b <+7>: pushl -0x4(%ecx)
0x0804853e <+10>: push   %ebp
0x0804853f <+11>: mov    %esp,%ebp
0x08048541 <+13>: push   %ecx
0x08048542 <+14>: sub    $0x4,%esp
0x08048545 <+17>: call   0x80484bb <passCheck>
0x08048548 <+22>: cmp    $0x1,%eax
0x0804854d <+25>: jne    0x8048556 <main+34>
0x0804854f <+27>: call   0x80484a3 <giveUserRoot>
0x08048554 <+32>: jmp    0x804855b <main+39>
0x08048556 <+34>: call   0x804848b <lockOutUser>
0x0804855b <+39>: sub    $0xc,%esp
0x0804855e <+42>: push   $0x0
0x08048560 <+44>: call   0x8048370 <exit@plt>
End of assembler dump.
```

CTF Link

- <http://goo.gl/forms/mivhCv3UELEgmDTM2>